



Ice Age melting down!
Intel features considered usefull!

XCon安全焦点信息安全技术峰会

XCon XFocus Information Security Conference

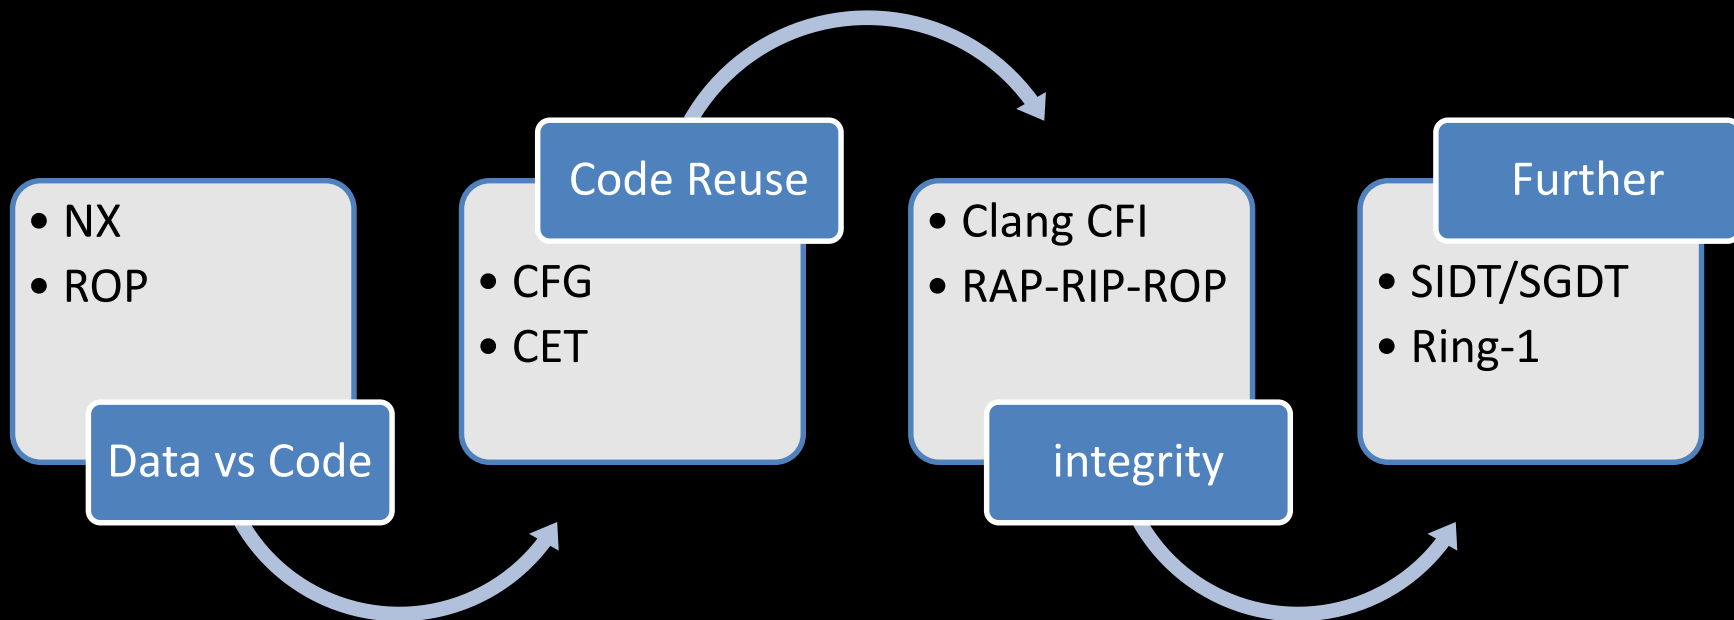


whoami

- @zer0mem, Peter Hlavaty
- Lead of Windows kernel security at KeenLab, Tencent
- 2015, 2016 pwn2own winner { team work! }
 - 2015 2x TTF
 - 2016 Master Of Pwn (Edge -> SYSTEM)
- Top 100 MSRC (Ranked 36 at 2016)
- Kernel CVEs and advanced exploitation techniques
- Delivered talks at Recon, SyScan, ZeroNights and others
- Wushu player



agenda






What is this talk about

- Elevation of Privilege
- Kernel code exec
- Software Mitigations
- New Bare metal principles
- Future of software security

15th

At the very begining

```
0D 03 0E B7 00 00 48 89 05 D1 31 03 00 C3 CC CC ...H8(F+.H8-(T...
48 83 EC 28 E8 7F BB 00 00 48 83 C4 28 E9 02 00 ...H8(F+.H8-(T...
00 00 CC CC 48 89 5C 24 10 48 89 74 24 18 57 48 ...H8(F+.H8-(T...
02 EC 20 E8 C0 00 00 00 0F D7 50 00 02 00 00 00 ...H8(F+.H8-(T...
```



you got code exec

```
; Attributes: library function
public start
start proc near
sub     rsp, 28h
call    sub_14003BB0C
add     rsp, 28h
jmp     __tmainCRTStartup
start endp
```

Custom data
shipped to
domain in
order to exec

- ExAllocatePool(...Nx)

ExAllocatePool allocates pool memory of the specified type and returns a pointer to the allocated block.

Syntax

C++

```
PVOID ExAllocatePool(  
    _In_ POOL_TYPE PoolType,  
    _In_ SIZE_T    NumberOfBytes  
);
```

Parameters

PoolType [in]

Specifies the type of pool memory to allocate. For a description of the available pool memory types, see [POOL_TYPE](#).

```
1 typedef enum _POOL_TYPE {  
2     NonPagedPool,  
3     NonPagedPoolExecute                = NonPagedPool,  
4     PagedPool,  
5     NonPagedPoolMustSucceed            = NonPagedPool + 2,  
6     DontUseThisType,  
7     NonPagedPoolCacheAligned            = NonPagedPool + 4,  
8     PagedPoolCacheAligned,  
9     NonPagedPoolCacheAlignedMustS      = NonPagedPool + 6,  
10    MaxPoolType,  
11    NonPagedPoolBase                    = 0,  
12    NonPagedPoolBaseMustSucceed         = NonPagedPoolBase + 2,  
13    NonPagedPoolBaseCacheAligned        = NonPagedPoolBase + 4,  
14    NonPagedPoolBaseCacheAlignedMustS  = NonPagedPoolBase + 6,  
15    NonPagedPoolSession                 = 32,  
16    PagedPoolSession                    = NonPagedPoolSession + 2,  
17    NonPagedPoolMustSucceedSession      = PagedPoolSession + 2,  
18    DontUseThisTypeSession              = NonPagedPoolMustSucceedSession + 2,  
19    NonPagedPoolCacheAlignedSession     = DontUseThisTypeSession + 2,  
20    PagedPoolCacheAlignedSession        = NonPagedPoolCacheAlignedSession + 2,  
21    NonPagedPoolCacheAlignedMustSSession = PagedPoolCacheAlignedSession + 2,  
22    NonPagedPoolNx                      = 512,  
23    NonPagedPoolNxCacheAligned           = NonPagedPoolNx + 4,  
24    NonPagedPoolSessionNx               = NonPagedPoolNx + 32,  
25 } POOL_TYPE;
```



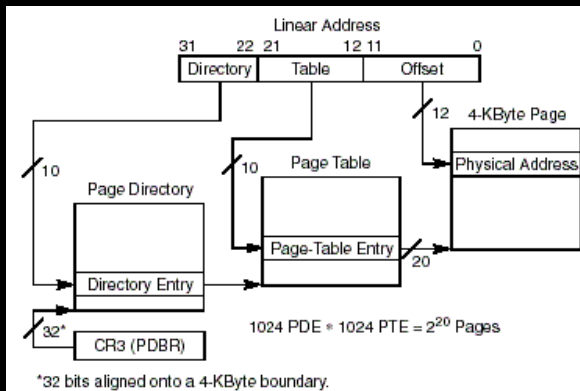
↓ : (

Custom data
shipped to
domain in
order to exec

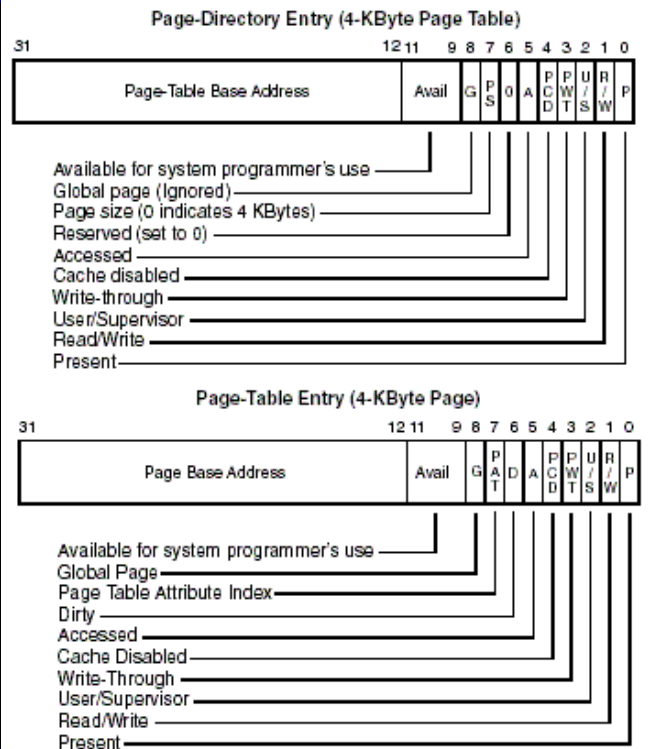
```
; Attributes: library function
public start
start proc near
sub     rsp, 28h
call    sub_14003BB0C
add     rsp, 28h
jmp     __tmainCRTStartup
start endp
```

Page Tables

- Exec bit on / off
- Accessible from ring0*




- * if no ring-1 in place



- Reusing of existing code
- Returning/Jumping into different (even misaligned) instructions
- Chaining those gadgets together
- Reliable control flow achieved

- Microsoft introduced mitigation
- Bitmap of valid indirect jump locations

- Simple
- Effective



<http://blog.trendmicro.com/trendlabs-security-intelligence/exploring-control-flow-guard-in-windows-10/>

- Check before indirect jump/call

- LdrpValidateUserCallTarget

```
mov     edx, dword ptr ds:GuardCFBitMapAddress
mov     eax, ecx
shr     eax, 8
mov     edx, [edx+eax*4]
mov     eax, ecx
shr     eax, 3
test    cl, 0Fh
jnz     short not_aligned_address
bt      edx, eax
jnb     short invalid_target
retn

not_aligned_address
or      eax, 1
bt      edx, eax
jnb     short invalid_target
retn
```

<http://powerofcommunity.net/poc2014/mj0011.pdf>

- Do not solve ROP itself, just focus on preventing kicking it off
 - fair enough approximation



Use function instead!

- pwn2own 2015
- Function driven attack
- Instead of gadget we use functions
 - Use it until we got custom kernel code exec



p2o 2015

- We got Read/Write primitive
 - ReCon , GdiBitmap technique
- Problems
 - #1 I dont like ROP and shellcode neither
 - #2 RWE page
 - #3 custom kernel code trigger

• NtUserMessageCall

```
extern "C"
void*
NtUserMessageCall(
    HWND hwnd, //window handle
    size_t fnSelector, //second arg - have to be bigger than 0x400
    size_t,
    size_t,
    size_t,
    size_t fnId //index-6 of function from MpFnIdPfn table
);
```

data	mpFnIdPfn	align	fn
.data:FFFFFFF97FFF3C0DB0	mpFnIdPfn	dq	?
.data:FFFFFFF97FFF3C0DB0			
.data:FFFFFFF97FFF3C0DB8	qword_FFFFFFFF97FFF3C0DB8	dq	?
.data:FFFFFFF97FFF3C0DC0	qword_FFFFFFFF97FFF3C0DC0	dq	?
.data:FFFFFFF97FFF3C0DC8	qword_FFFFFFFF97FFF3C0DC8	dq	?
.data:FFFFFFF97FFF3C0DD0	qword_FFFFFFFF97FFF3C0DD0	dq	?
.data:FFFFFFF97FFF3C0DD8	qword_FFFFFFFF97FFF3C0DD8	dq	?
.data:FFFFFFF97FFF3C0DE0	qword_FFFFFFFF97FFF3C0DE0	dq	?
.data:FFFFFFF97FFF3C0DE8		db	? ;
.data:FFFFFFF97FFF3C0DE9		db	? ;
.data:FFFFFFF97FFF3C0DEA		db	? ;
.data:FFFFFFF97FFF3C0DEB		db	? ;
.data:FFFFFFF97FFF3C0DEC		db	? ;
.data:FFFFFFF97FFF3C0DED		db	? ;
.data:FFFFFFF97FFF3C0DEE		db	? ;
.data:FFFFFFF97FFF3C0DEF		db	? ;
.data:FFFFFFF97FFF3C0DF0		db	? ;
.data:FFFFFFF97FFF3C0DF1		db	? ;
.data:FFFFFFF97FFF3C0DF2		db	? ;

<http://www.k33nteam.org/noks.html>

15th

p2o 2015

ExAllocatePool

Provide RWE
flags

Leak address



p2o 2015

- Use `gdi.io()` to write our kernel mode driver in place
 - `cc_shellcode` framework with some kernel mode features
- Need to deal with relocations, and import resolving
 - `cc_shellcode` framework do it automatically



p2o 2015

- Kernel code exec switch
 - gdi.io() _EPROCESS list walking
 - Find stack of our thread
 - With gdi.io() rewrite own stack return



CET - IBT

- At compile time mark source and destination of call
 - Indirect call/jmp
 - ENDBRXX
- If you jump somewhere what is not marked as destination, boom
- Similar to CFG
 - roots better into the code
 - this time from bare-metal setting distinctions
- Direct and clean solving of ROP
 - Solve ROP comprehensively not just at very beginning!
 - Basically shaping architecture forward to secure oriented one
- Does *not* intend to solve integrity!

<https://software.intel.com/sites/default/files/managed/4d/2a/control-flow-enforcement-technology-preview.pdf>



Code Exec

- anti-ROP is no problem
- No direct page-table altering is no problem
- Type-hash integrity + stack, ouch...
- EPT ? ...



EPT & Code Signing

- Running at ring-1
- No W+X pages at once, no ever
- EPT vs ring0 page tables, shadowing
- Signing check once it tries to exec
- Game over for custom code exec if no logical bug ?



Control Flow control

- CodeSigning or PageTable lockdown no problem
 - We dont intend to use custom code
- anti-ROP becomes a problem +-
 - Type-hash integrity is essential problem



Control Flow Integrity

Another technologies
[not applied at windows kernel]

XCon安全焦点信息安全技术峰会

XCon XFocus Information Security Conference



RAP-RIP-ROP

- Return address protection
- Reserved register for cookie
- Cookie per :
 - Task
 - Syscall
 - Special loops scenarios

<https://pax.grsecurity.net/docs/PaXTeam-H2HC15-RAP-RIP-ROP.pdf>

RAP-RIP-ROP

```
push %rbx
mov 8(%rsp),%rbx
xor %r12,%rbx
...
xor %r12,%rbx
cmp %rbx,8(%rsp)
jnz .error
pop %rbx
retn
.error:
ud2
```

- cookie ^ ret
- Kernel.io() from another thread can be possibly used to leak & rewrite
 - However unread-able kernel stacks can harden approach

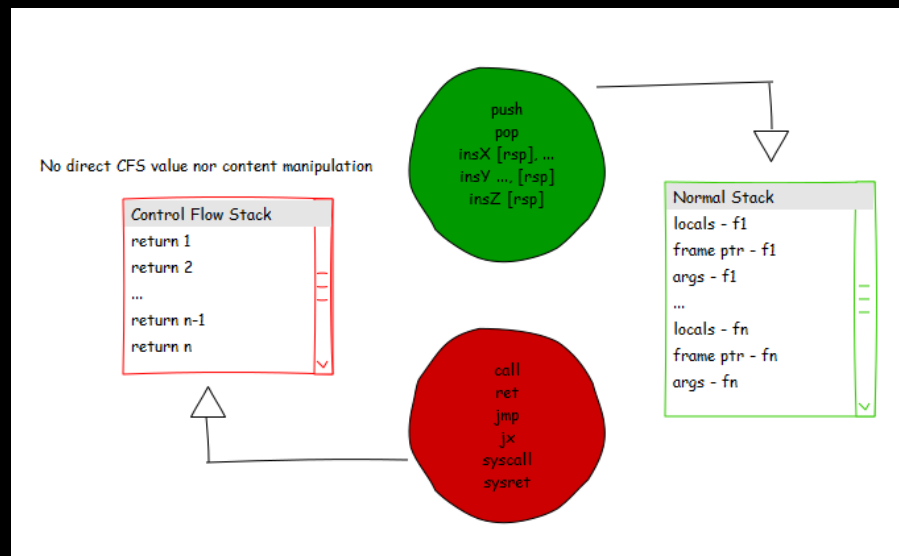
Self-modify

- Limited / jailed kernel-A.io() to create better version of kernel-X.io()
- kernel-X.io() need :
 - Contains loop in which do read/write
 - Loop counter should depend on writable memory
- iovec alike are good candidates (splice, ...)

Self-modify

- kernel-X.io() modify itself
- Corrupt loop counter
- To create prolonged loop
- Read out own cookie^ret
- Rewrite own cookie^ret
 - In another thread resolve and modify
 - Readed value in previous step must be accessible somehow (user mode / kernel.io()) to second thread
 - Patched value should be available in further read/write operations of self-looped kernel-X.io()

- Normal instruction can not touch it directly
- Kernel.lo() is out of game





Clang

- Forward-Edge CFI
- bit vector for static types
 - Read-only
 - Per static type
- Compiler + linker
 - Depends on LLVM's type metadata

<http://www.pcc.me.uk/~peter/acad/usenix14.pdf>

```

struct A {
    virtual void f1();
    virtual void f2();
    virtual void f3();
};

struct B : A {
    virtual void f1();
    virtual void f2();
    virtual void f3();
};

struct C : A {
    virtual void f1();
    virtual void f2();
    virtual void f3();
};

```

The scheme will cause the virtual tables for A, B and C to be laid out consecutively:

Virtual Table Layout for A, B, C														
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
A::offset-to-top	&A::rtti	&A::f1	&A::f2	&A::f3	B::offset-to-top	&B::rtti	&B::f1	&B::f2	&B::f3	C::offset-to-top	&C::rtti	&C::f1	&C::f2	&C::f3

The bit vector for static types A, B and C will look like this:

Bit Vectors for A, B, C															
Class	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
A	0	0	1	0	0	0	0	1	0	0	0	0	1	0	0
B	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0
C	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0

<http://clang.llvm.org/docs/ControlFlowIntegrityDesign.html>



RAP-RIP-ROP

- Indirect Control Transfer protection
- Type hash for indirect 'calls'
 - Return type
 - Function name
 - Function parameters

RAP-RIP-ROP

```
cmpq $0x11223344,-8(%rax)
jne .error
call *%rax
...
cmpq $0x55667788,-16(%rax)
jne .error
call *%rax
...
dq 0x55667788,0x11223344
func:
```



Perfect ?

- Still imperfect, however infinite times better!
- Exponentially raise cost (time) of exploitation
- Security becoming real deal



Intel – errata

- Some may argue that intel technology is weak because dont do type-hash etc
- However goal of this is not provide CFI, but do strict rule how to jump (kill ROP)
- Simple & effective
- CFI need to be handled compiler specific



Bit of c++

- Is this correct ?
- Is it even important ?
- Could it ends up with control flow control ?

```
class A
{
    virtual void foo(x,y,z);
}
class B : A
{
    virtual void foo(x,y,z);
}
..
B b;
...//meanwhile is arbitrary rewritten b.foo = &A.foo
b.foo(x,y,z);
```



Bit of c++

- Well ... yes and no ...
 - depends on quality of code
 - Inheriting vs limiting / extending, etc
 - concrete object context
 - Must be used with (limited) Kernel.Io()
 - Could allows to reach data gadget dispatcher
 - ... what ?

- Turing complete

```
1 //dispatcher & jump :
2 void cmd_loop(server_rec *server,conn_t *c) {
3     while (TRUE) {
4         pr_netio_telnet_gets(buf, ...);
5         cmd = make_ftp_cmd(buf, ...);
6         pr_cmd_dispatch(cmd); //calls functions
7         // with memory errors and gadgets
8     }
9 }
10 char *pr_netio_telnet_gets(char *buf,...) {
11     while(*pbuf->current != '\n' && toread>0)
12         //reads through virtual PC
13         *buf++ = *pbuf->current++;
14 }
```

Code 10. Gadget dispatcher and simulated jump gadget. pbuf->current is the virtual PC pointing to the malicious input.

- (Conditional) jump operation. Code 10 shows the ProFTPD program logic to read the next command from an input buffer. pbuf->current is a pointer to the next command in the input, thus forming a virtual PC for the attacker's MINDOP program. By corrupting pbuf->current, the attacker can select a particular input that invokes a specific MINDOP operation. We use the assignment operation to conditionally update the virtual PC, thus simulating a conditional jump operation.

https://www.comp.nus.edu.sg/~shweta24/publications/dop_oakland16.pdf

Data Oriented Prg

- However type-hash approaches ..
- #1 You have luck to hit vulnerability reachable at control flow of some gadget dispatcher
- #2 you need to find type-hash friendly candidate and prepare for redirection to it before you proceed to DOP



Bit of objective-C

```
void foo(x,y,z);  
void fxx(x,y,z);  
  
struct A  
{  
    decltype(&foo)* AFoo;  
}  
..  
A a;  
a->AFoo = foo;  
...//meanwhile is arbitrary rewritten a.AFoo = &fxx  
a.AFoo(x,y,z);
```

- This is more powerful to abuse
- some languages are build to be dynamic ..
some of them not
 - Therefore harder to protect

OOP problems

- Data – Code mix
 - Objective-C
 - Hard to impossible to distinguish
 - C++
 - Can be distinguished
 - Built for dynamic code in mind

Joe Armstrong [↗] (2011)

Why OO Sucks [↗]

"Objects bind functions and data structures together in indivisible units. I think this is a fundamental error since functions and data structures belong in totally different worlds."



Data flow integrity

- DFI – data flow integrity
- Different set of instructions are usually responsible for some set of data
- Likely to work on same type of data



DFI - implementation

- Static data flow graph at compilation
- Checking per memory access ?
- Costly overhead
- Approximations ?



DFI - approximation

- Limit only to subset of functions
- Limit only to subset of data
- Leaves lot of attack surface
- Approximate more ?



DFI - approximation

- Maybe help of hardware ?
- Special instruction :
 - change *accessible range* mem-access instructions
 - Default should be current stack
 - Whoaa, somehow close to intel's mpx! ☺
- Fine grained pools & inside isolated heaps alike
- Compilers + linkers with type scope information
- Data attacks will be solved 一步一步

<https://software.intel.com/en-us/blogs/2013/07/22/intel-memory-protection-extensions-intel-mpx-support-in-the-gnu-toolchain>



DFI - approximation

- However .. Based on current (pretty badass) implementation of DOP and its early results
- Seems to be fair enough to just protect potential gadget dispatchers
 - Numbers are tend to be small
- But this is more like state-of-art, and subsets of this approach still remains (kernel.io()!)

- Those are results of non - security thinking (limitations) at the very beginning
- Seems to be solved
 - Unless logic bugs in solutions
 - However idea of solutions is pretty solid



Code Reuse & Data attack

- Can we say with all those potential mitigations (hardware & software) will be solved ?

Code Reuse & Data attack

- Not likely
- As those exploiting very principles of current architectures & dynamic code
 - Proper protection is hard (but looks like possible) for built-in dynamic languages
 - To what extend is it possible to languages dynamic by possibility not be design ?

Security

- However ... with ongoing research (clang, pax, intel, ..)
- It will be far more harder to do successful attack
 - in many (many) cases will be even impossible
- Goal of success are however not solely mitigations, nor bugs killing itself
 - Both aspects is important and has its own role
 - Final outcome could be very secure environments



Back to the future

Still way to go..

XCon安全焦点信息安全技术峰会

XCon XFocus Information Security Conference




Time to think

- What is SYSTEM / Kernel Code Exec ?
- Why we do need it ?
- How is it used ?
- What is real goal we want to achieve ?

15th

DATA



10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101
10010101010101010010101010101



- DATA are all what matters
- Code is just group of targeted instructions to work over those data
- You want read / write to data
- SYSTEM / kernel code exec did it for you in easy way



Kernel.io()

- We have it from the time being
- Kernel.io() is subset of DOP implementation
- We used it for making shortcuts
- Once you can read / write to targeted domain you won
 - no need to direct control flow control



Kernel.io()

- Kernel.io()
 - Vulnerability
 - Considering mem-corruption here
 - Technique
 - Leaking domain addresses
 - Reliable re-use multiple times
 - Deep domain knowledge
 - Code quality++


```
bool
Write(
    __in_bcount(size) void* addr,
    __in_bcount(size) const void* buff,
    __in_size_t size
) override
{
    return Io<true>(addr, const_cast<void*>(buff), size);
}

bool
Read(
    __in_bcount(size) const void* addr,
    __inout_bcount(size) void* buff,
    __in_size_t size
) override
{
    return Io<false>(const_cast<void*>(addr), buff, size);
}
```



Kernel.io()

- with Kernel.io() you can :
 - do any mathematic operation on arbitrary data
 - emulate known 'api'
 - With no protection of stack-return it is pretty cheap




security

- SIDT / SGDT
- CR4.UMIP (documented at revision 58)
- Solves (trap to block) :
 - Leaks for user mode
 - Virtualization leak issue

- gdi info leakage
locked



Nicolas Economou @NicoEconomou 07 Jul
GDI objects memory leak (PEB, GdiSharedHandleTable) is fixed in "Windows 10-Insider Preview 14342"..
killed vector :(



```
kernel debug (x86)
kd> dq poi (win32kbase!gpcntHmgr)
ffff8322'80210000 00000000'00000000 00000000'00000000
ffff8322'80210010 00000000'00000000 00000000'00000000
ffff8322'80210020 00000000'00000000 00000000'00000000
ffff8322'80210030 00000000'00000000 00000000'00000000
ffff8322'80210040 00000000'00000000 00000000'00000000
ffff8322'80210050 00000000'00000000 00000000'00000000
ffff8322'80210060 00000000'00000000 00000000'00000000
ffff8322'80210070 00000000'00000000 00000000'00000000
kd> d
ffff8322'80210080 00000000'00000000 00000000'00000000
ffff8322'80210090 00000000'00000000 00000000'00000000
ffff8322'802100a0 00000000'00000000 00000000'00000000
ffff8322'802100b0 00000000'00000000 00000000'00000000
ffff8322'802100c0 00000000'00000000 00000000'00000000
ffff8322'802100d0 00000000'00000000 00000000'00000000
ffff8322'802100e0 00000000'00000000 00000000'00000000
ffff8322'802100f0 ffffffff'ff00000a 00040004'00000000
kd>
ffff8322'80210100 00000000'00000000 ffffffff'ff00000b
ffff8322'80210110 00080088'00000000 00000000'00000000
ffff8322'80210120 ffffffff'ff00000c 00080008'00000000
ffff8322'80210130 00000000'00000000 ffffffff'ff00000d
ffff8322'80210140 00080008'00000000 00000000'00000000
ffff8322'80210150 ffffffff'ff00000e 00080008'00000000
ffff8322'80210160 00000000'00000000 ffffffff'ff00000f
ffff8322'80210170 00050085'00000000 00000000'00000000
kd>
```



Security [WINDOWS]

- w32k lockdown
 - dont touch what you dont need
- ntos restricted access
 - Low attack surface
 - Code quality good

Conclusions

- Hardware setting security boundaries at possibility level
- Software continues to tackling attackers techniques to the edge
- Vulnerabilities space is shrinking
 - Notable by security research from china



Way to go



PaX Team

@paxteam

@redragonvn @daveaitel i'll be the first one to switch to it if it works as return address protection is the slower part of RAP.

- Technologies benefits one from another, making software more secure!

15th

Thank you!
Q & A

科恩实验室
KEEN
security lab

