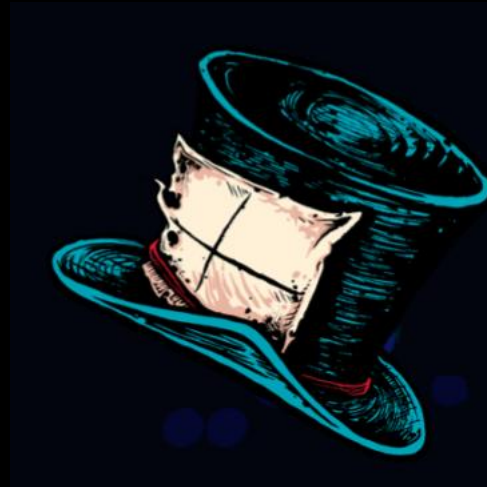


You didnt see it's coming?

“Dawn of hardened Windows Kernel”



\$whoami

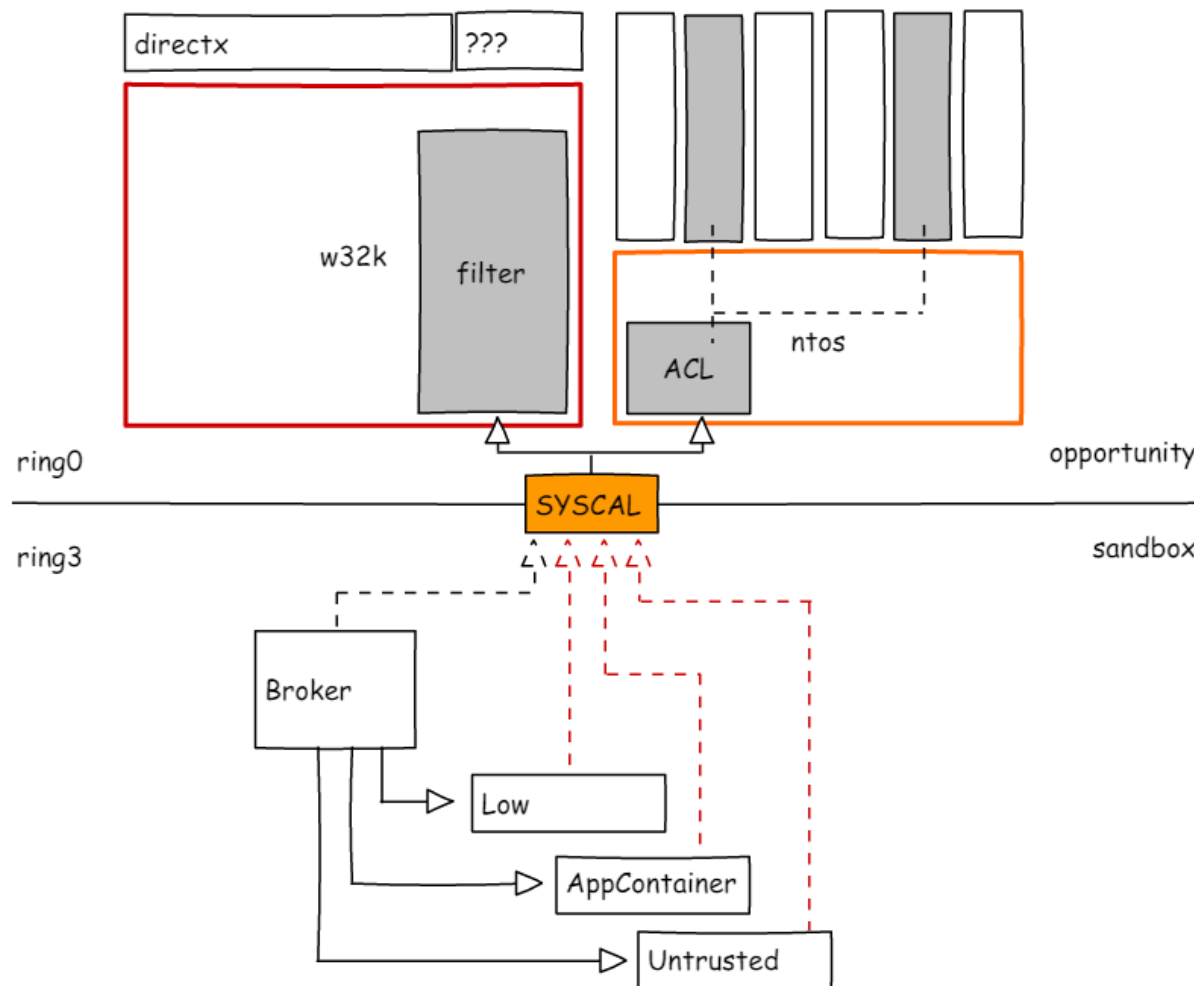
Peter

- @zer0mem
- Windows kernel research at KeenLab, Tencent
- pwn2own winner (2015 / 2016), pwnie nominee (2015)
- fuzzing focus : state
- wushu player

agenda

- Kernel attack surface
- Fuzzing
- Exploitation techniques
- Mitigations

windows sandbox kernel attack surface



path of less resistance - w32k*

- resides in ring 0
 - natural := (level < medium) -> SYSTEM (/kernel code exec) escape
- huge attack surface
 - huge in comparsion to ntoskrnl counterpart or in-ring3-sandbox interface
 - state logic, window callbacks, hidden syscalls, directx, format parsing, ...
- accessible from sandbox-es
 - nowadays more or less => big success!
- easy develepoing of exploitation techniques

w32k vs 2016

MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-017
MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-017
MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-017
MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-017
MS16-062	Win32k Information Disclosure Vulnerability	CVE-2016-017
MS16-062	Microsoft DirectX Graphics Kernel Subsystem Elevation of Privilege Vulnerability	CVE-2016-017
MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-019
MS16-062	Win32k Elevation of Privilege Vulnerability	CVE-2016-019

win32k 1 of 33 Options ▾

Security Advisories and Bulletins > Acknowledgments ▾

2016

Acknowledgments – 2016

MS16-090	Win32k Elevation of Privilege Vulnerability	CVE-2016-3249
MS16-090	Win32k Elevation of Privilege Vulnerability	CVE-2016-3250
MS16-090	GDI Component Information Disclosure Vulnerability	CVE-2016-3251
MS16-090	Win32k Elevation of Privilege Vulnerability	CVE-2016-3252
MS16-090	Win32k Elevation of Privilege Vulnerability	CVE-2016-3254
MS16-090	Microsoft win32k Elevation of Privilege Vulnerability	CVE-2016-3286

Win32k Elevation of Privilege Vulnerability	CVE-2016-3308
Win32k Elevation of Privilege Vulnerability	CVE-2016-3309
Win32k Elevation of Privilege Vulnerability	CVE-2016-3310
Win32k Elevation of Privilege Vulnerability	CVE-2016-3311

what is going on ?

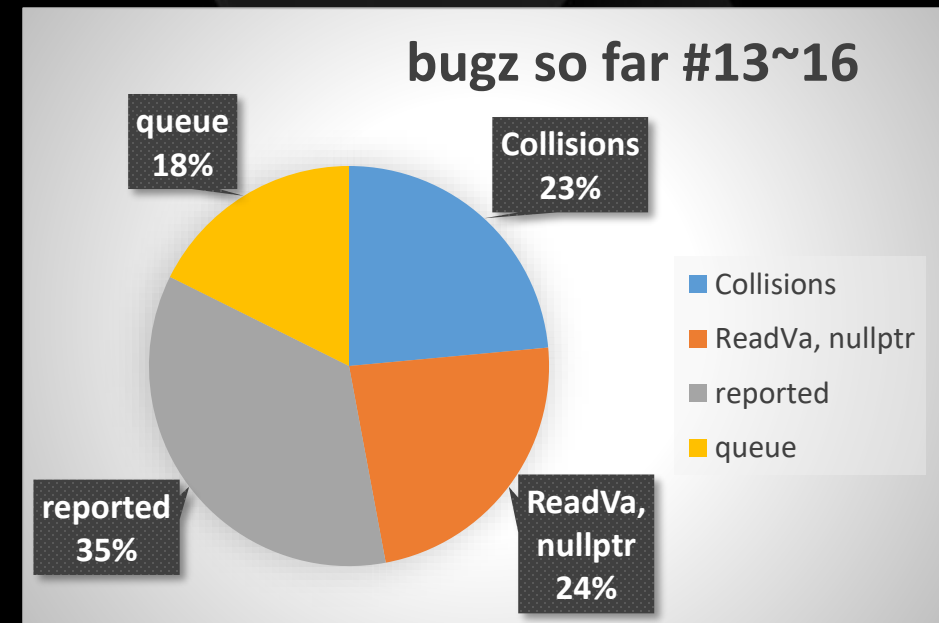
- huge numbers of syscalls
- lot of objects
- lot of hardcore graphics stuffs
- lot of things i dunno



```
C:\>cat w32k@subsurface | grep "Nt" | wc -l  
1042
```

w32k vs Qilin

- Qilin - internal multi-platform fuzzing framework
- Target – gdi - part of w32k
- Technique :
 - knowledge based
 - random driven
 - interconnection aware
- Results in graph
- technique ++ code coverage
 - feedback + knowledge based (ongoing)

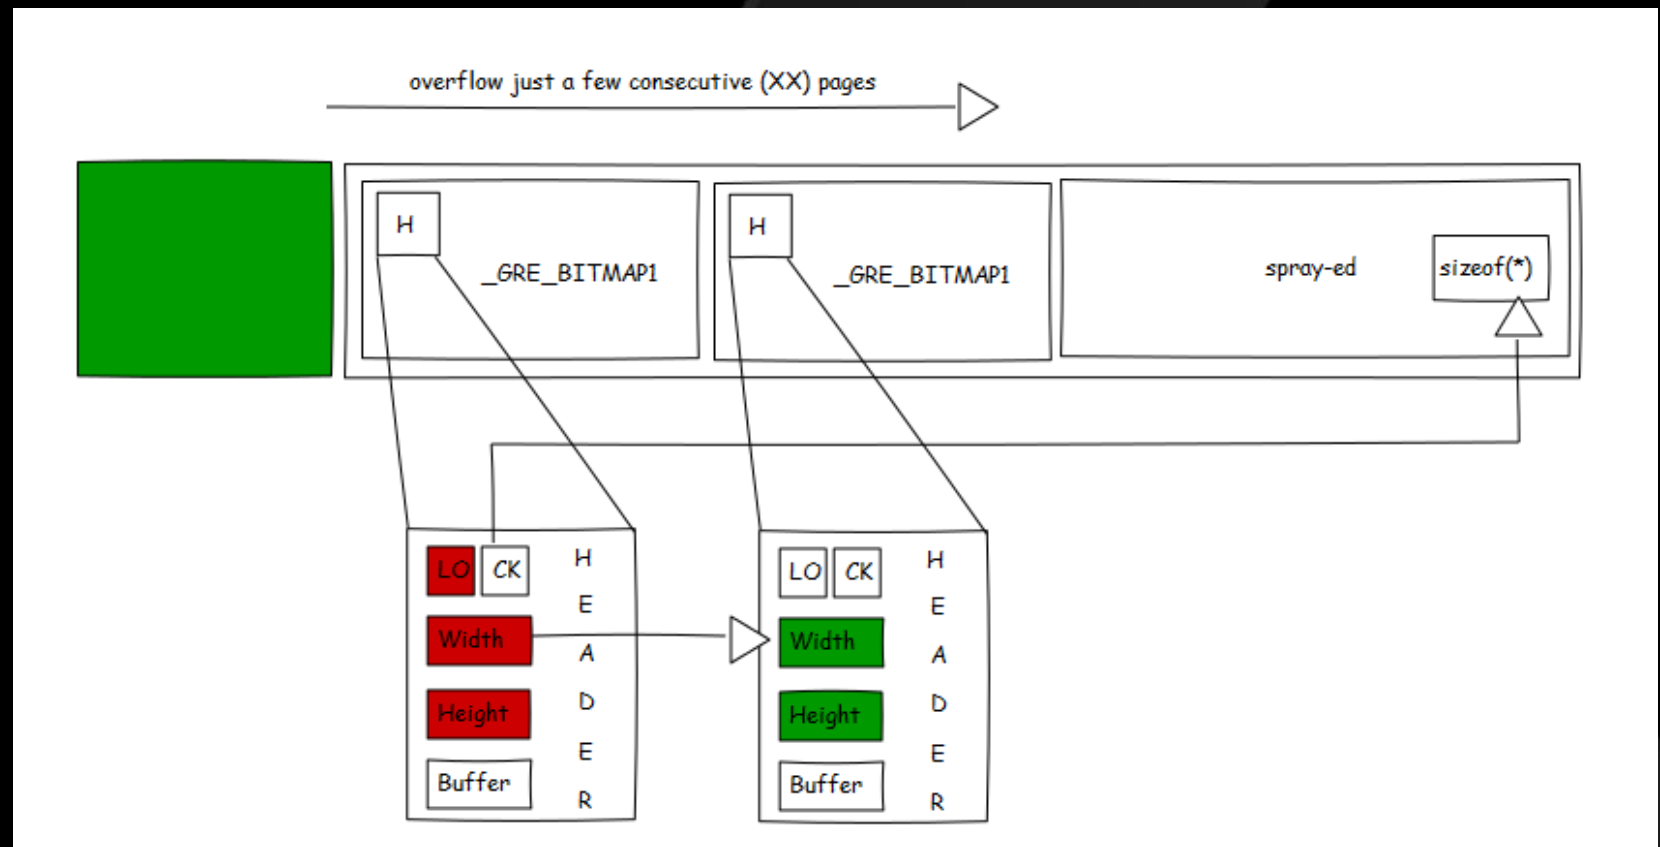


from bug to take over SYSTEM

- except huge code base with lot of space for bugs
- huge code base with lot of space for missusing existing mechanisms
- valuable arsenal
 - virtual tables and alike
 - unprotected lengths
 - plain state members
 - buffer pointers
- prerequisites :
 - call syscalls

W32k kernel io : _gre_bitmap

- Simple & *reliable* pool layout
- Direct syscalls for content manipulation
- Pivot -> worker technique



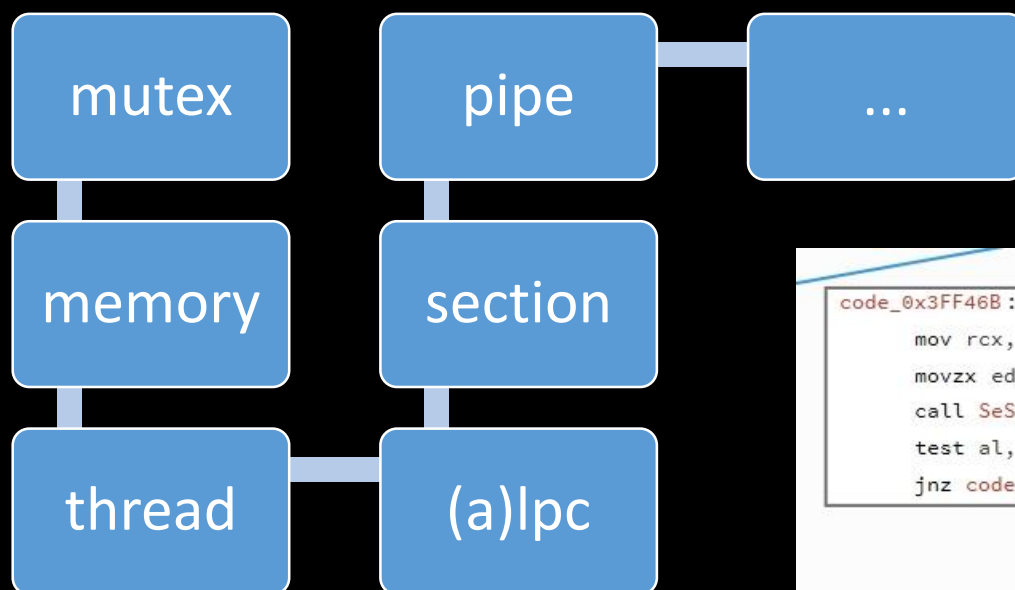
w32k out of scope of interest ?

- somehow :
 - going to be locked down & filtered out
 - not so cool anymore, many research done, many bugs out, ...
 - so complex so it even gets boring after time of researching there
 - if you trying to find something -> you will : if no, then try harder...
- however in nowadays still applies (and some time will)
 - weakest points of windows kernel
 - accessible from most sandboxes
 - best attack vector nowadays
 - attackers are lazy, why do something hard if no need ?

PERFECT GROUND FOR YOUR FUZZER



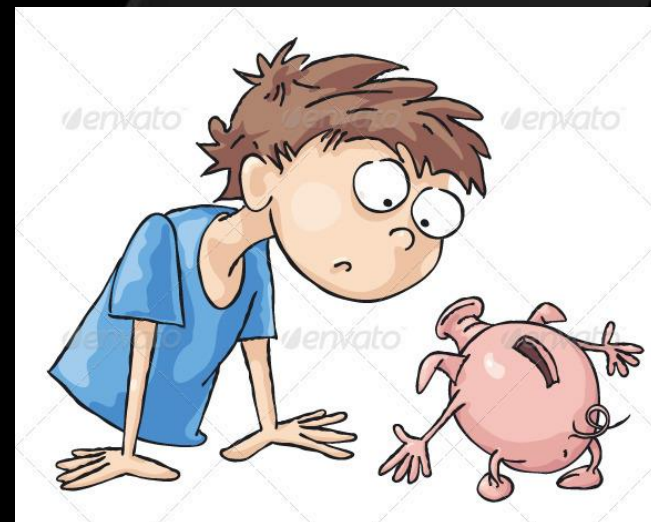
ntos : lets move from win32k



```
code_0x3FF46B :  
    mov rcx, qword [data_0x742CE0]  
    movzx edx, r12b  
    call SeSinglePrivilegeCheck ; unsigned char __cdecl SeSinglePrivilegeCheck( ur  
    test al, al  
    jnz code_0x3FF3C6  
  
code_0x3FF483 :  
    mov edi, 0xC0000061
```

ntos - attack surface

- difference :
 - seemingly no data parsing +/-
 - state only
 - from untrusted level / app container not much to touch
 - small number of state changing syscalls
 - state changes are minimalistic in most cases*



ntos under the microscope - extensions

- Nt*Transaction*
- Nt*Enlistment*
- Nt*Manager*

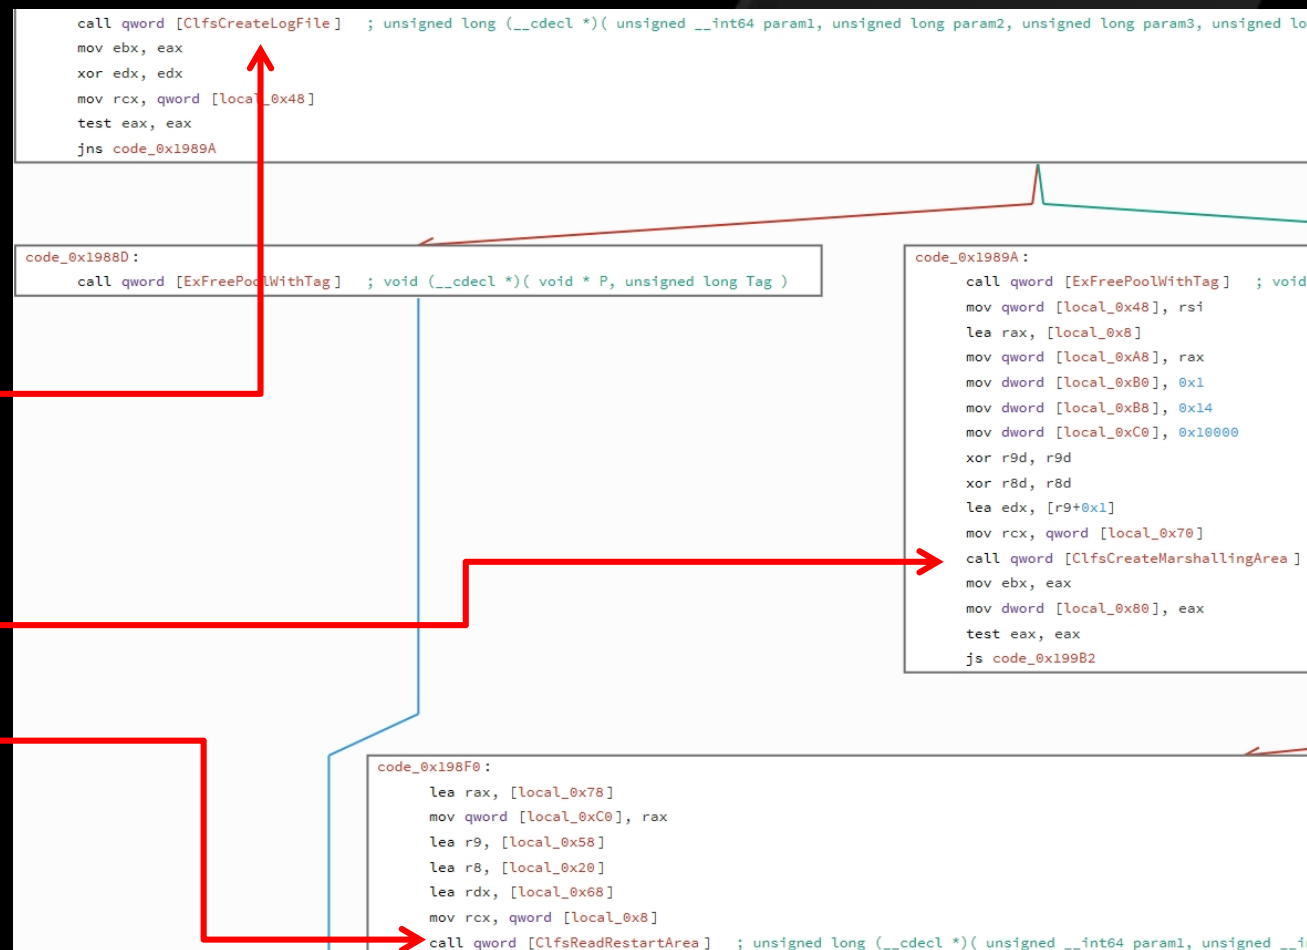
Strings Imports Functions			
Type	Length	Value ▲	
Import	8	NtCommitTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtCreateTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtCreateTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtEnumerateTransactionObject	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtFreezeTransactions	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtOpenTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtOpenTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtQueryInformationTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtQueryInformationTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtRecoverTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtRenameTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtRollbackTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtRollforwardTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtSetInformationTransaction	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtSetInformationTransactionManager	ext-ms-win-ntos-tm-l1-1-0.dll
Import	8	NtThawTransactions	ext-ms-win-ntos-tm-l1-1-0.dll

tm.sys

- ntos extension
 - Kernel Transaction Manager
 - however not much attack surface as well
 - simple states & minimum syscalls
 - however interconnections
-
- Qilin vs tm, round1 :
 - 1 nullptr
 - 1 exploitable double ObDeref -> type confusion -> uaf

explore unknown : indirections++

- tm.sys simple purpose driver
- but interesting module involved at backend
- CLFS.sys



CLFS.sys : unseen hand behind the scenes

- Common Log File System
- not everything need to be direct
- easy to get interest if you look at simplistic tm.sys
- clfs used at many different places in ntos itself as well
- clfs on other hand more complex
 - c++ code base
 - complex state
 - involve data parsing
- Qilin vs clfs : 3~5 state bugs, 6~8 data parsing
- Advanced *data* fuzzer : covers all first hand bugs, and extend 4+
 - by @long123king

bug hunting

research, fuzzing, feedback, analysis, loop back

kick off

- select target
- collect related syscalls (apis)
- research attack surface
 - go through related msdn part (if any)
 - understand abstracted functionality
 - what is its purpose
 - what is usual way of working with it
 - double check interesting points with kernel implementation

fuzzer

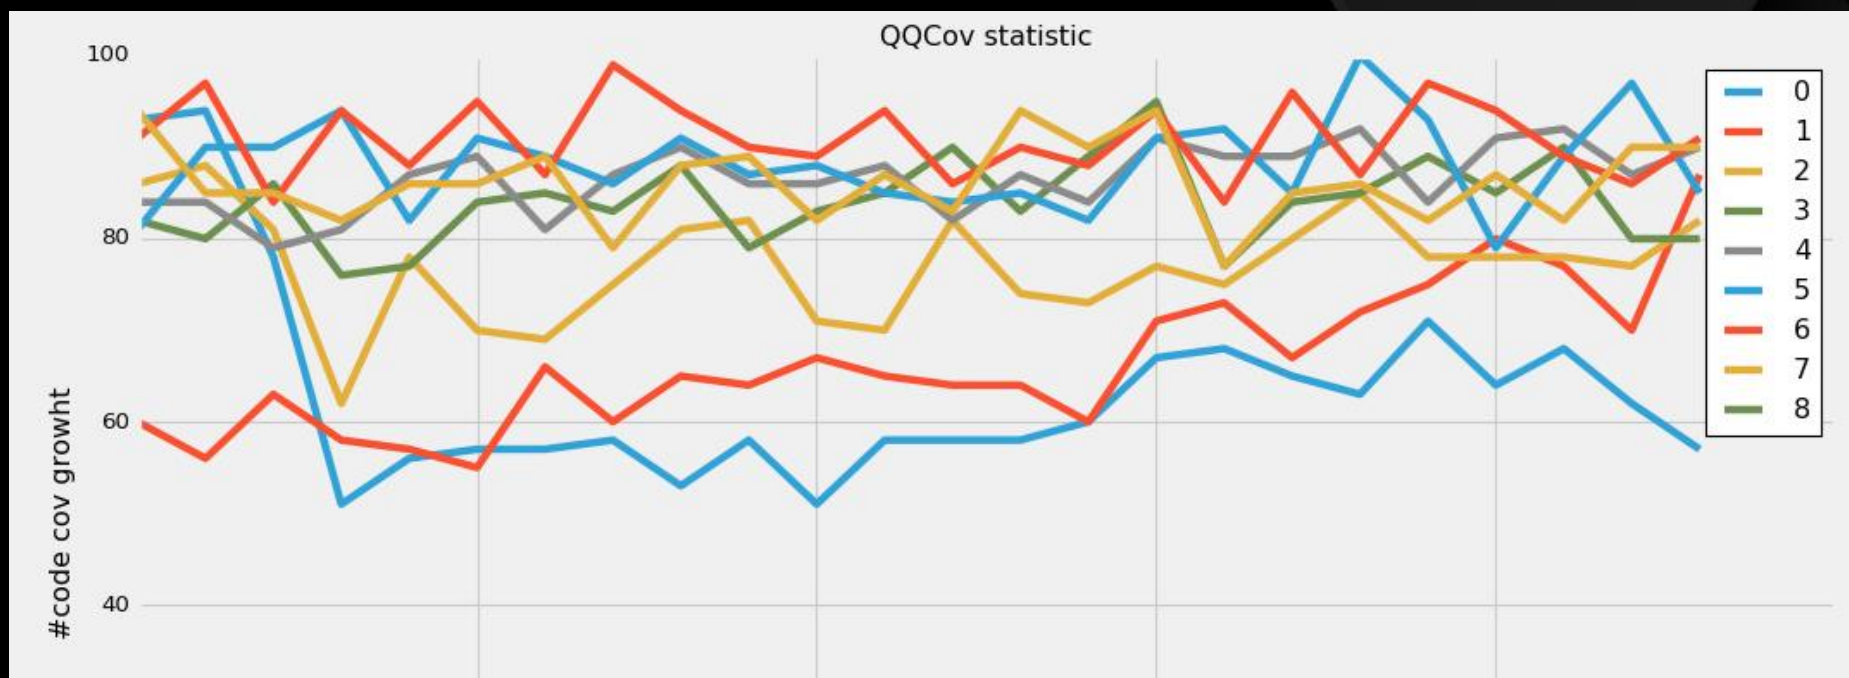
- make some code to automate working with target
 - optionally use existing framework, just add functionality per target
- make it work
 - most basic syscalls working (open, close, get)
- make use of previous reseach
 - implement basic schema from previous step
 - ensure that your fuzzer can sucessfully work on target based on documentation of target
 - introduce randomnes with preserving certain sucess ratio
 - introduce connections (syscall & handle dependencies)

feedback

- how good is your fuzzing ?
- success ratio % ?
- code coverage ?

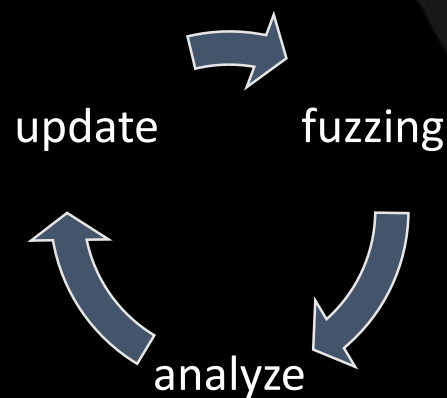
<pre>#overall elapsed time : [1:24] Total Process Count : 1 Process Killed : 2 Total Server # : 0 Total Client # : 0</pre>	<pre>#ring0 Total KObj count : 5 Dropped KObj # : 3046 Total Syscalls # : 6165 Syscalls # : 976 [15%] Average Syscall Total # : 55</pre>
<pre>#kobjs server-not-found server-not-found server-not-found server-not-found server-not-found</pre>	
<pre>#syscalls ConDrvConLockedOr => total : 129, ratio : 100.00 NtWriteFile => total : 195, ratio : 100.00 CdWriteIoOutput => total : 252, ratio : 100.00 CdpLaunchServerProcess => total : 99, ratio : 100.00 ConHostAPI => total : 165, ratio : 73.33 ConDrvConFastIoctl => total : 39, ratio : 100.00 NtCreateFile => total : 1023, ratio : 0.00</pre>	

measure your fuzzing



research

- fuzzing is one part
 - you want more bugz ? you need know your target well
 - not only target, but your fuzzer as well
-
- analyze less covered parts
 - think of possible buggy scenarios
 - implement back to fuzzer
 - loop back (from documentation, trough fuzzing up to understanding)



additional thoughts

- fuzzing is not enough
 - you need to understand what you are fuzzing
 - better understanding helps to build more tricky logic
 - however once you concrete too much, then you putting yourself into corner
- better to be backuped with code coverage corpus
 - corpus is essential to keep & update, but not recycle (scatter&replace) often
- fuzzer is not all you need
 - runtime tools [address sanitizers, race tools, debugger plugins, ..]
 - static tools [helping your fuzzer, helping your analysis]

lets move to exploitation
techniques!

options

- w32k easy to go, however not so easy from locked down / filtered state anymore
 - can go through chaining, but we can do better
- ntos, harder to go
 - not much available syscalls - harder to setup & control technique properly
 - not much used before - why to go harder way if no need ?
 - less objects, with less nasty states

nt!KeWaitForSingleObject

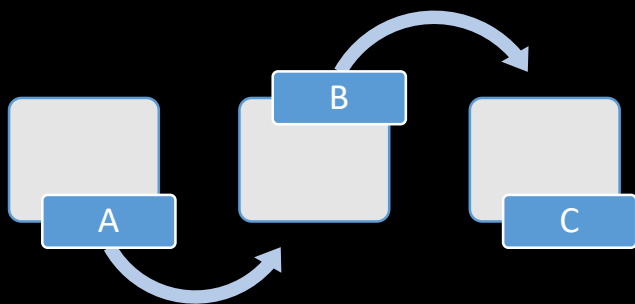
- small object
- relatively simple
- interesting logic behind usage
 - lock / unlock
 - wait
 - kernel pointers
 - ethreads
- offers various primitives :
 - arbitrary decrease
 - arbitrary write where - kernel pointer*
 - ...

```

nop     dword ptr [rax+rax]
inc     ebp
test    dword ptr [nt!Hv1LongSpinCountMask (fffff802`ce0281c0)],ebp
je      nt! ?? ::FNODOBFM::`string'+0x39f0 (fffff802`cdddbb0)
pause
mov     eax,dword ptr [rdi]
test    al,al
js      nt!KeWaitForSingleObject+0x440 (fffff802`cdcabc610)
lock bts dword ptr [rdi],7
jb      nt!KeWaitForSingleObject+0x440 (fffff802`cdcabc610)
lea     rdx,[rbx+140h]
```

SafeLink write where primitive

- Introduced to deal
meta data link corruptions
- previously lead to :
write where - what
- now it leads to int 0x29
... not necessary, or ?

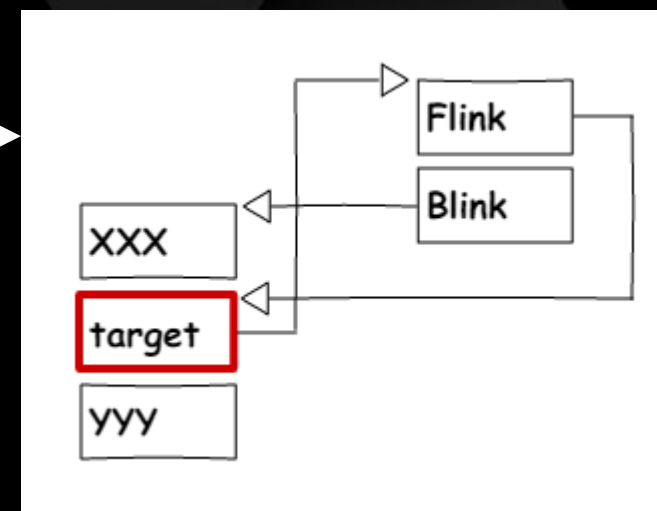
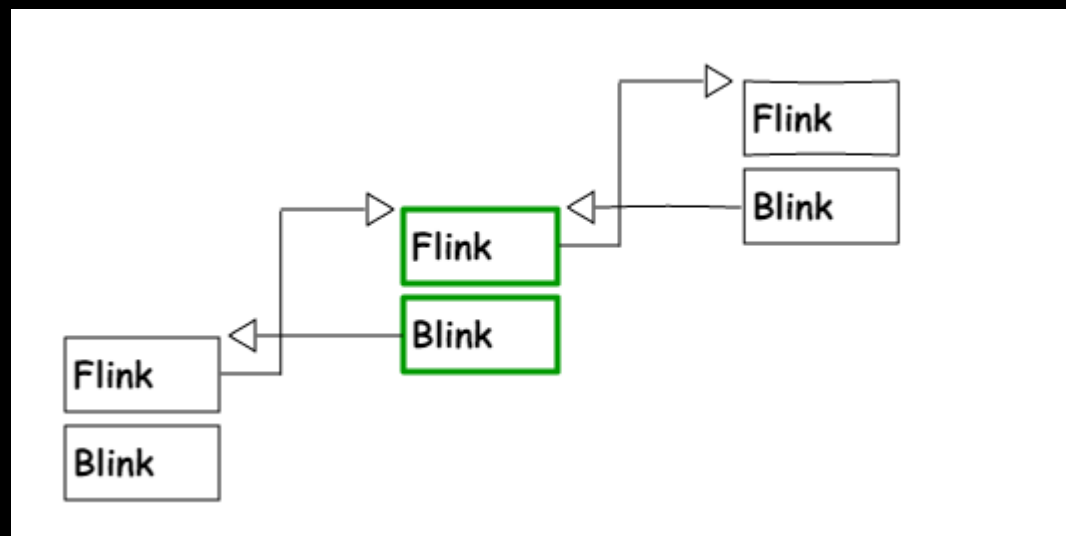


```
code_0x1D2B3:
    ebp, 0x102
    code_0x1D079:
        mov rcx, qword [r14+0x10]
        lea rax, [r14+0x8]
        cmp qword [rcx], rax
        jnz code_0x1D28D
    code_0x1D28D:
        mov ecx, 0x3
        int 0x29
    code_0x1D294:
        movzx ecx, byte
        movzx eax, cl
        and al, 0x7
        cmp al, 0x1
    code_0x1D08A:
        mov qword [r10], rax
        mov qword [r10+0x8], rcx
        mov qword [rcx], r10
        mov qword [rax+0x8], r10
```

Arbitrary write – kernel pointer

- target needs to point to semaphore / event
- Problem ? .. Well not much :
 - You can predict kernel mode memory members
 - You can misalign

Flink -> Blink == Blink -> Blink



ALPC

- not everything touchable in ntos is really minimalistic
- more complex mechanism
- more complex state
- good interconnections
- .. however well written code ..
 - lot of safe checks
 - safe user mode memory handling (via 'getters')
 - seems well designed

ALPC #spray

```
kd> dt nt!_kalpc_message
```

```
...
```

```
+0x010 PortQueue
```

```
+0x018 OwnerPort
```

```
...
```

```
+0x068 MessageAttributes
```

```
+0x0b0 DataUserVa : Ptr64 Void
```

```
...
```

```
+0x0d8 ExtensionBuffer : Ptr64 Void
```

```
+0x0e0 ExtensionBufferSize : UInt8B
```

```
+0x0e8 PortMessage : _PORT_MESSAGE
```

```
kd> dt nt!_kalpc_message PortMessage->u1
```

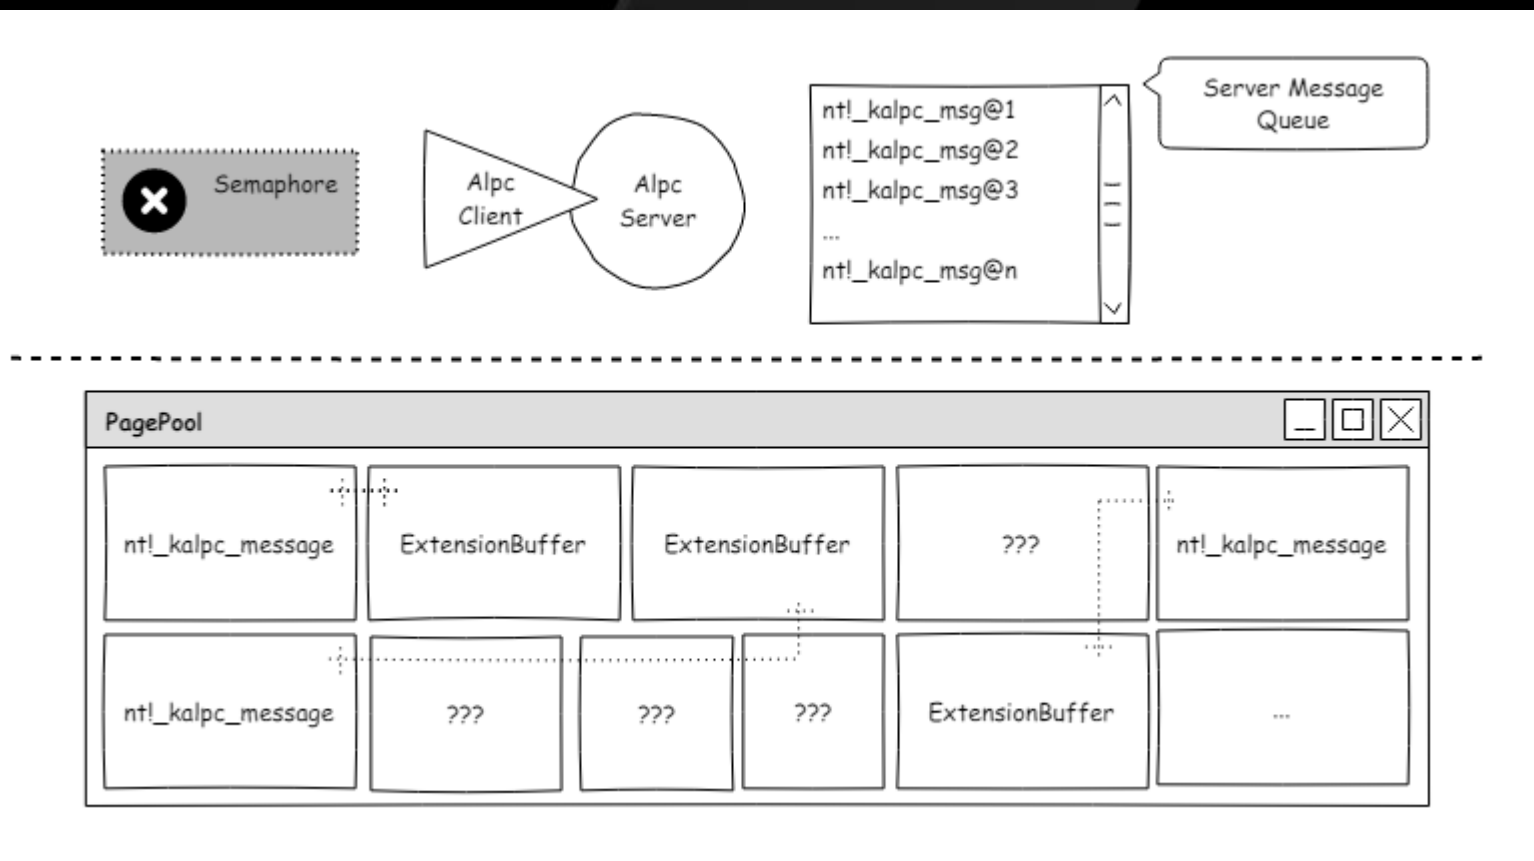
```
+0x0e8 PortMessage :
```

```
+0x000 u1 :
```

```
+0x000 s1 : <unnamed-tag>
```

```
+0x000 Length : UInt4B
```

```
+0x004 u2
```



ALPC #detect

```
kd> dt nt!_kalpc_message
```

```
...
```

```
+0x010 PortQueue
```

```
+0x018 OwnerPort
```

```
...
```

```
+0x068 MessageAttributes
```

```
+0x0b0 DataUserVa : Ptr64 Void
```

```
...
```

```
+0x0d8 ExtensionBuffer : Ptr64 Void
```

```
+0x0e0 ExtensionBufferSize : UInt8B
```

```
+0x0e8 PortMessage : _PORT_MESSAGE
```

```
kd> dt nt!_kalpc_message PortMessage->u1
```

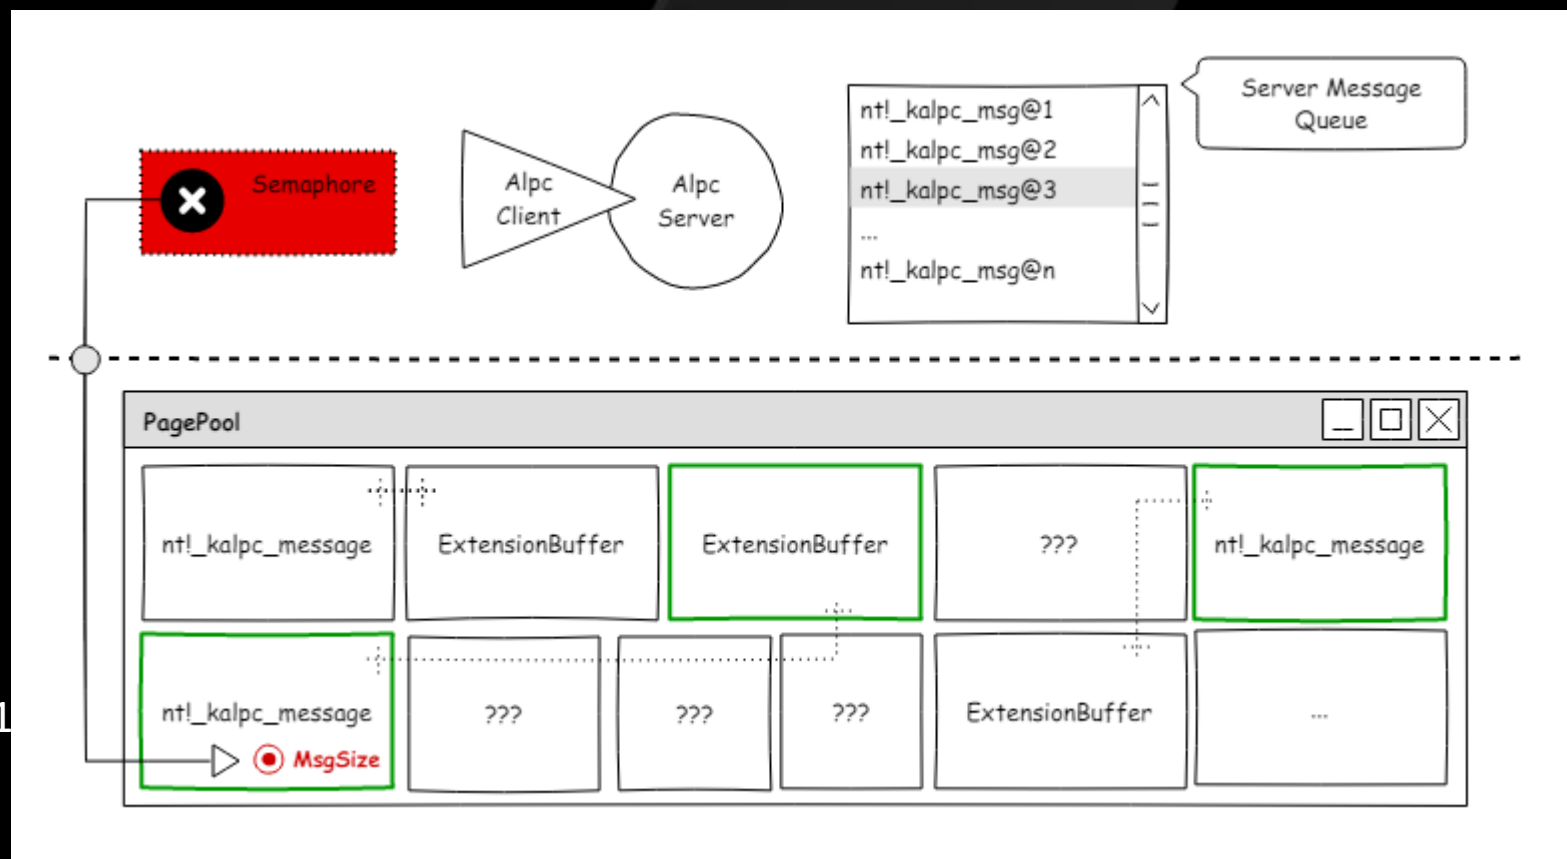
```
+0x0e8 PortMessage :
```

```
+0x000 u1 :
```

```
+0x000 s1 : <unnamed-tag>
```

```
+0x000 Length : UInt4B
```

```
+0x004 u2
```



ALPC #overflow

```
kd> dt nt!_kalpc_message
```

```
...
```

```
+0x010 PortQueue
```

```
+0x018 OwnerPort
```

```
...
```

```
+0x068 MessageAttributes
```

```
+0x0b0 DataUserVa : Ptr64 Void
```

```
...
```

```
+0x0d8 ExtensionBuffer : Ptr64 Void
```

```
+0x0e0 ExtensionBufferSize : Uint8B
```

```
+0x0e8 PortMessage : _PORT_MESSAGE
```

```
kd> dt nt!_kalpc_message PortMessage->u1
```

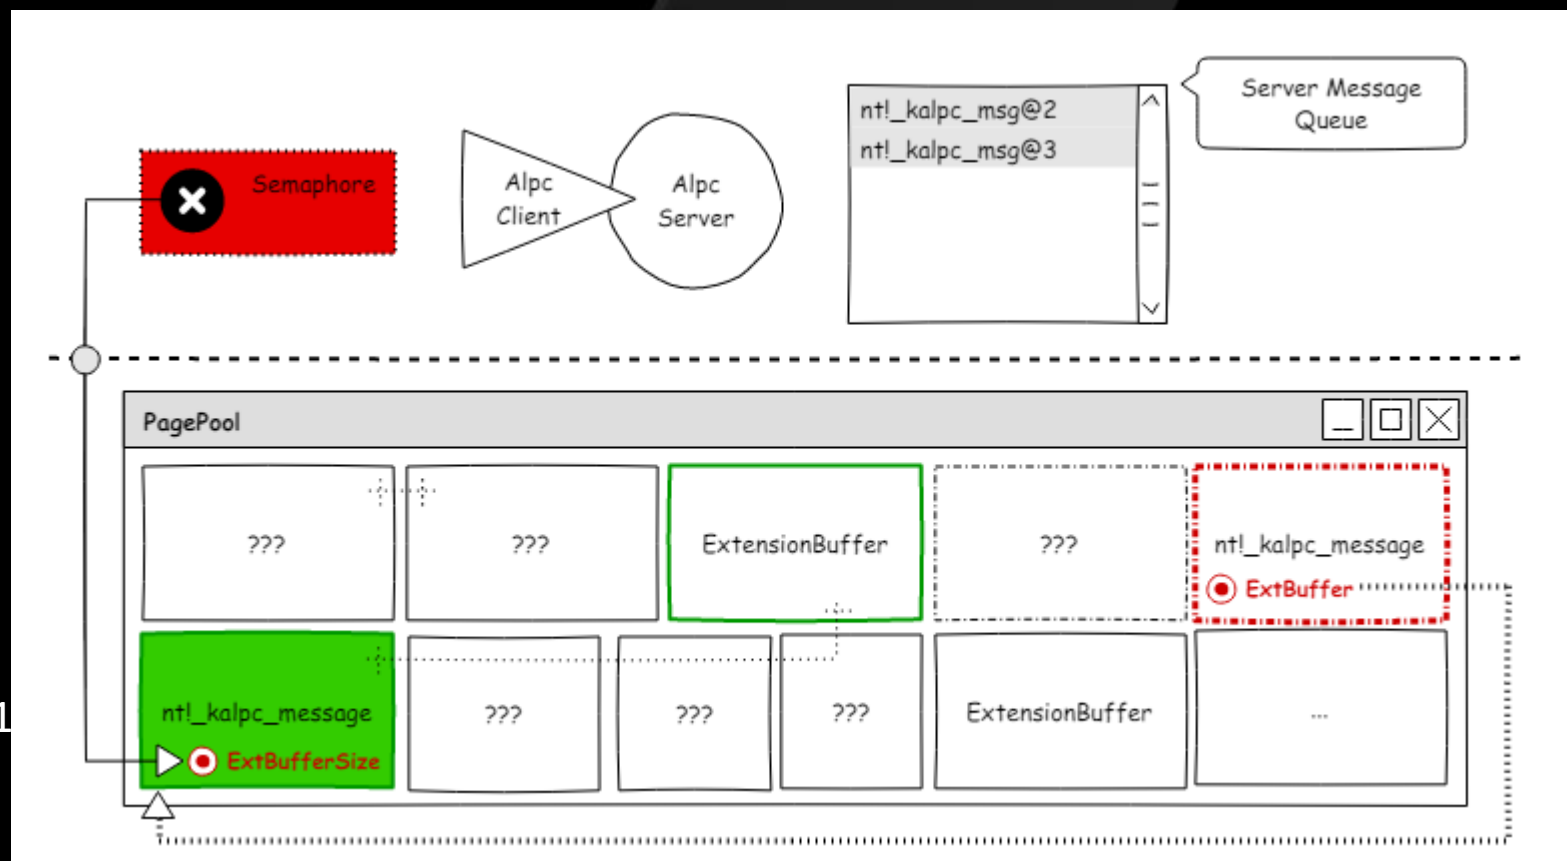
```
+0x0e8 PortMessage :
```

```
+0x000 u1 :
```

```
+0x000 s1 : <unnamed-tag>
```

```
+0x000 Length : Uint4B
```

```
+0x004 u2
```



ALPC #io

- #4 threads
(gdi tech, only #1 thread)

- io
- pivot
- worker
- server

```
void
PushRequest(
    CAlpcMsg& msg
)
{
    CAlpcClientWorker worker(m_client, msg);
    worker.SendTcp();//let server to handle request
}
```

- Performing additional write per io !
- synchronization
- Undocumented ALPC
 - Setup server – client
 - Keeping messages alive
- pool feng shui is problematic

```
void
KernelIoWrite(
    uint8_t* addr,
    const uint8_t* mem,
    size_t size
)
{
    if (!size)
        return;
    if (size % sizeof(void*))
        return;

    m_ioAddr = nullptr;

    std::thread pivot(PushRequest, this, m_pivot);

    while (!m_ioAddr)
        Sleep(100);

    std::thread worker(PushRequest, this, m_worker);

    Sleep(200);//wait until worker will be in queue

    memcpy(m_worker.IoData(), mem, sizeof(void*));

    m_ioAddr = addr; // pivot start operate!

    pivot.join();
    worker.join();

    return KernelIoWrite(
        addr + sizeof(void*),
        mem + sizeof(void*),
        size - sizeof(void*));
}
```

Hardenings !?

SMEP, KASLR, Nx, CFG, RFG, ACL, SafeLink ...

from user to kernel, from kernel to jail : w32k

- Filtering

- ? win32k – win32kfull – win32kbase ?
 - win32k -> wrapper to add stub* with access checks
- restrict access only to necessary parts
 - maybe not 'clean' solution, but security++ (relatively)
- app can have different sandboxed entities, and per entity different attack surfaces
 - find way from one to another, and you can find way out of filter
- in optimal scenario limit potential bug landscape
 - first part of good effect
- in optimal scenario limit exploitation techniques
 - this is the second good effect

- lockdown

- wow, this sure solve something in crazy *absolute* way O_o

w32k is here for you if you need it!

- One can have arbitrary decrease ? [see previous semaphore slides]
 - Or any suitable primitive, just necessary be creative
- Would he choose alpc technique ?
- If he is lazy (smart) enough, then no ...

```
ffff9784`e97a4800 1020 MicrosoftEdge.  
ffff9784`e874c800 1764 browser_broker  
ffff9784`cdbbc800 1824 MicrosoftEdgeC  
0: kd> dt _eprocess ffff9784`cdbbc800 EnableFilteredWin32kAPIs DisallowWin32kSystemCalls AuditFilteredWin32kAPIs  
nt!_EPROCESS  
+0x300 DisallowWin32kSystemCalls : 0y0  
+0x6c4 EnableFilteredWin32kAPIs : 0y1  
+0x6c4 AuditFilteredWin32kAPIs : 0y1
```

- Re-enable w32k instead!
 - even allow more proc in your job in case of need
- do one-bit kernel pwn via win32k!_gre_bitmap and enjoy life!

kernel code exec vs mitigations

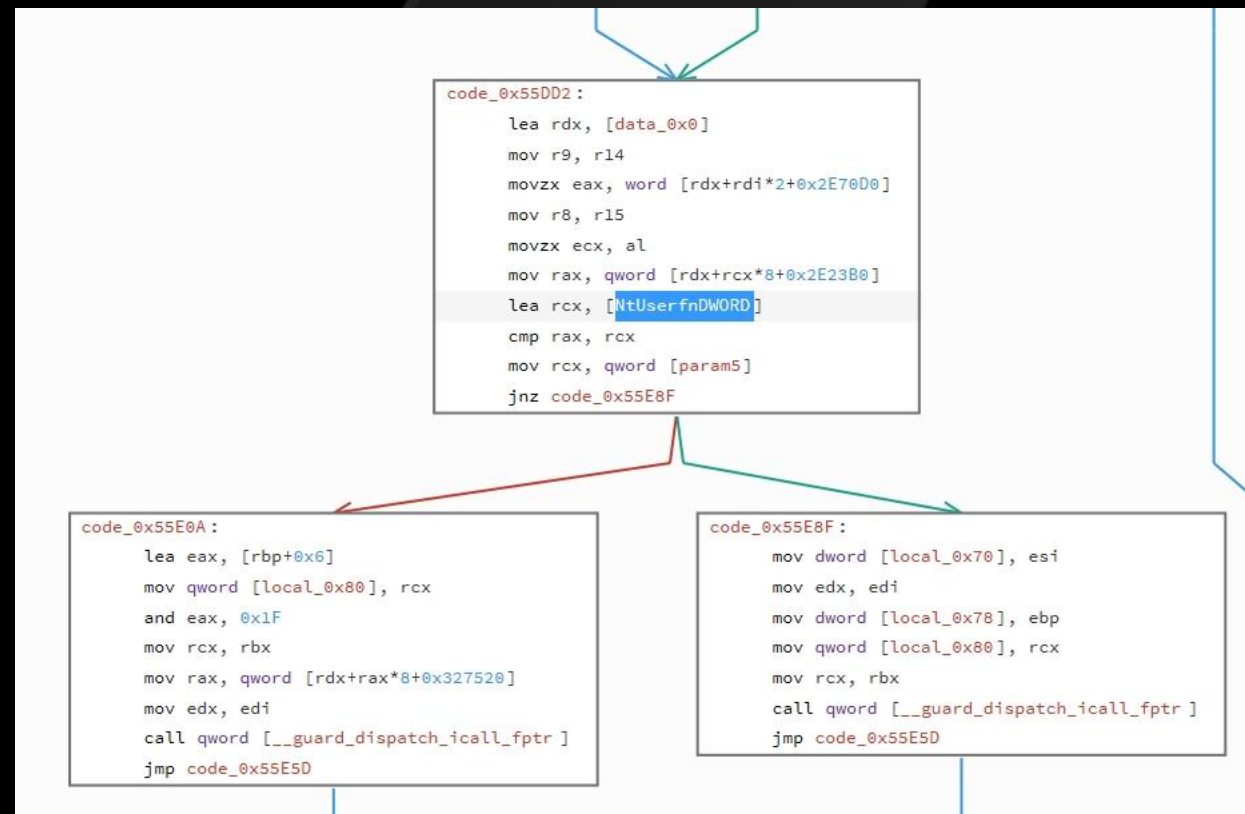
- KASLR, code signing, Nx, CFG, SMEP
- set some good security boundaries
- however getting code exec no problem after all
 - kaslr - good bug or additional info leak
 - code signing : we will ship our code to kernel via kernel io, and exec it
 - Nx + SMEP : find rwe page, or create it
 - CFG : find good trampoline or target stack

easy to be said, harder to be done ?

- not actually
- kernel-io techniques already described
- still headache with kernel Nx / SMEP ?
 - NtUserMessageCall [following slide]
- CFG
 - prevents only from ROP to be kicked of
 - i dont like ROP neither, target stack [+ functions]
 - CFG alone is not enough (in terms of code exec) !

NtUserMessageCall

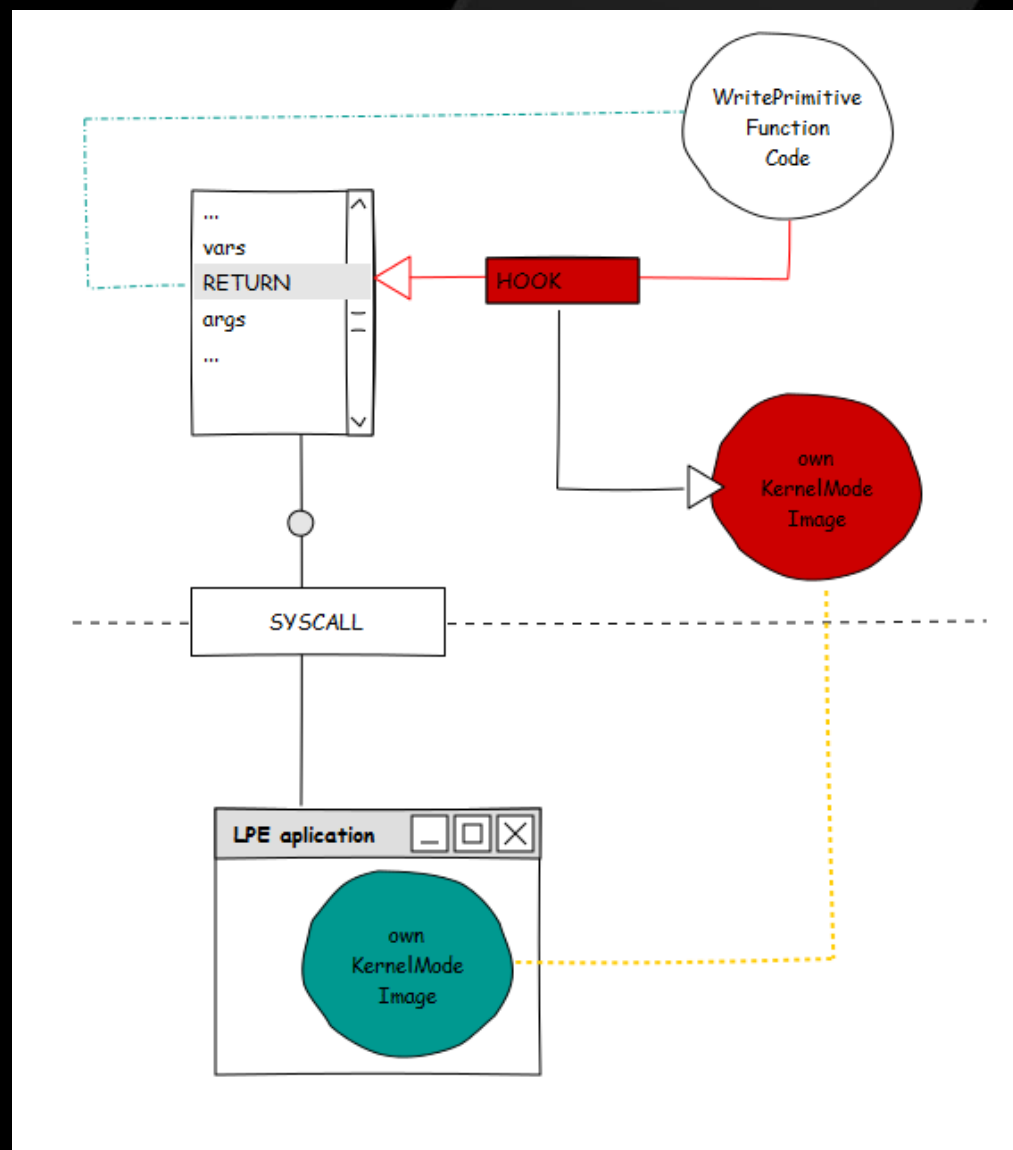
- Function of many faces
- 7 parameters
- Full control over them
- Indirect call invoked
- 7+ member of vtable <- dummy!
- vcall return is passed to user!



- Enough primitives for kernel code exec, with valid function (no ROP) !

Get kernel code exec

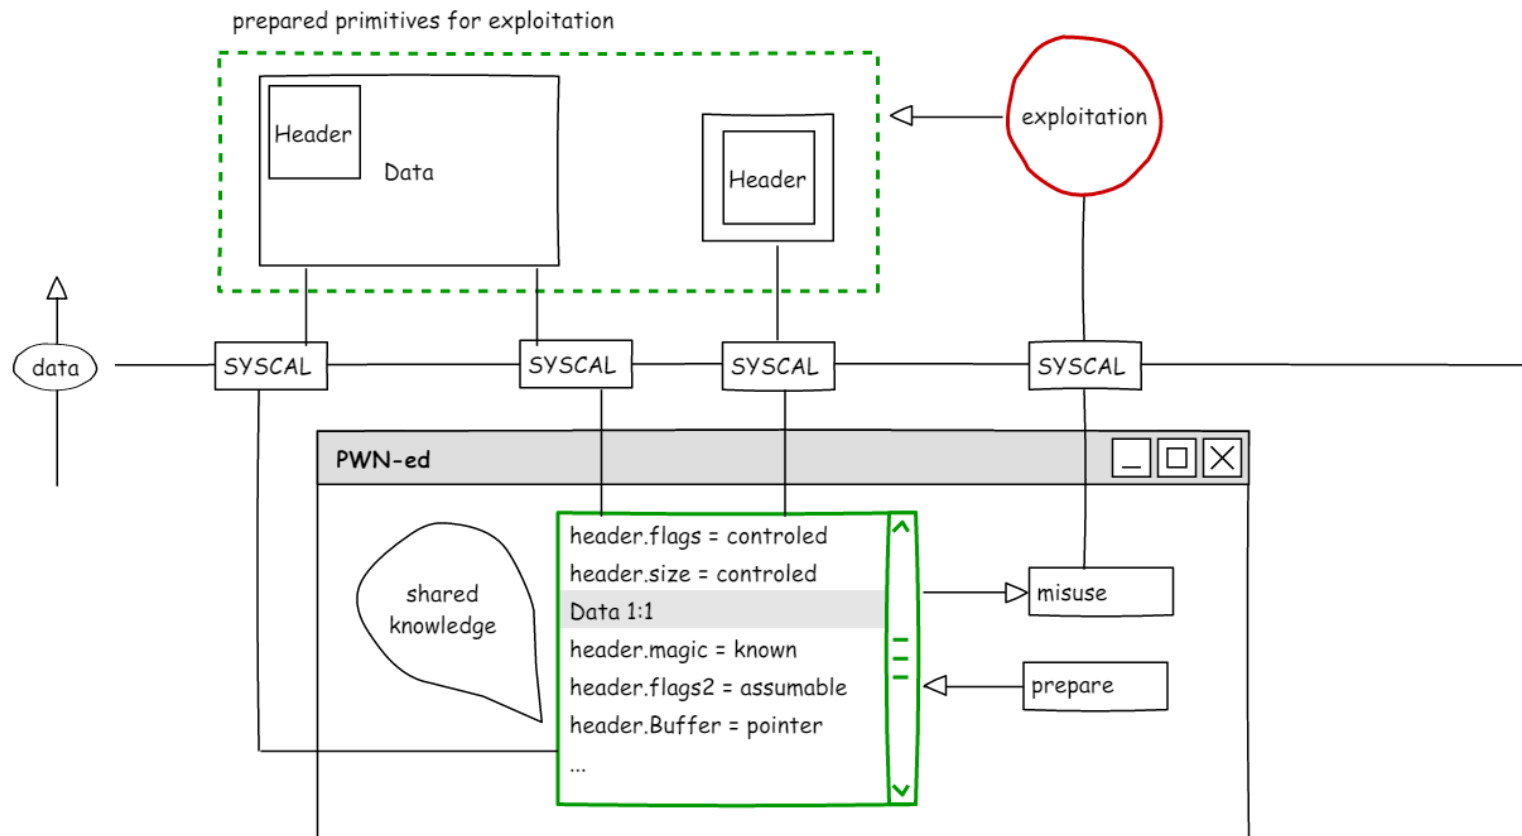
- Misuse existing functions
 - NtUserMessageCall
 - Allocate RWE memory
 - memcpy you kernel driver
 - Stack hook
-
- Game over



SMAP - tale of lacking feature

- prevents kernel unwanted access to user mode memory
- sets bare metal borders between data in user mode and kernel
- should serve as fine security feature ... does it ?

SMAP - tale of lacking feature



SMAP - tale of lacking feature

... however ...

- without proper security schema in underlying OS, it is just shinny ...
 - nullptr deref protection - anything bellow 0x10000
 - poisons should be **not-mappable** by user by default!
- you need controlled data at exploitation ? why you need them in user mode ?
- easy to put in kernel (pipes, direct mem mapping, physmap, kernel stack, ...)
- in occasional cases one more bug to leak address layout
 - in most cases you can get them out from mem corruption bug (unprotected raw pointers, ...)
- security features is good to have, but better to use to max potential
 - which OS nowadays is using SMAP in way that it is really obstacle for attackers ?

Control Flow – hardware!

- CET-IBT
 - Indirect jumps / calls + endbrXX
 - Similar idea with CFG indeed
 - Your jumps can not go wild, like ROP do!
- CET-Shadow Stack
 - nicely done!
 - You can not subvert control flow via stack hooking anymore!
- so seriously, what now ?
 - ... wait for some time to adapt to mainstream O_o

Return Flow Guard

- Rumors about RFG
 - Return Flow Guard
 - Preliminary documentation (nice!) :
<http://xlab.tencent.com/en/2016/11/02/return-flow-guard/>
 - Simple & fast & effective
- Effective ?
 - Shadow Stack alike, therefore yes!
 - .. and we are back with FS on x64! fs:[rsp]
 - **arbitrary read / write will not touch it!**
 - in combination with CFG ? – hell, a lot!
 - Two simple & effective mitigations together are pretty solid!
 - Can be implemented with todays hardware!

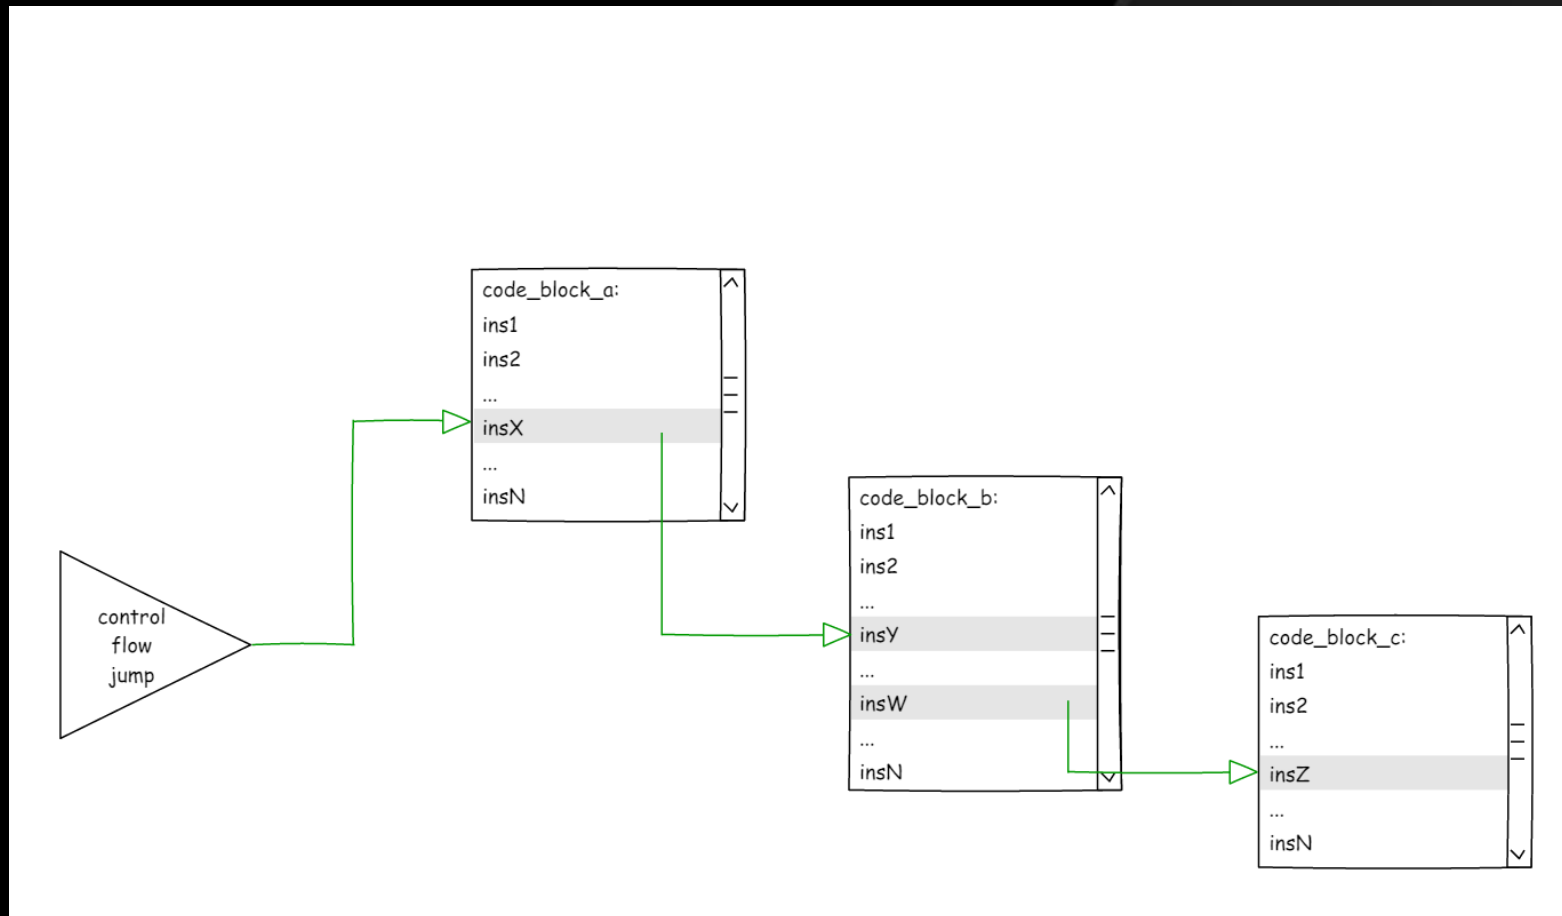
3.4 RUNTIME INSERTED EPILOGUE BYTES (15 BYTES)

```
MiRfgInstrumentedEpilogueBytes
mov    r11, fs:[rsp]
cmp    r11, [rsp]
jnz    _guard_ss_verify_failure
```

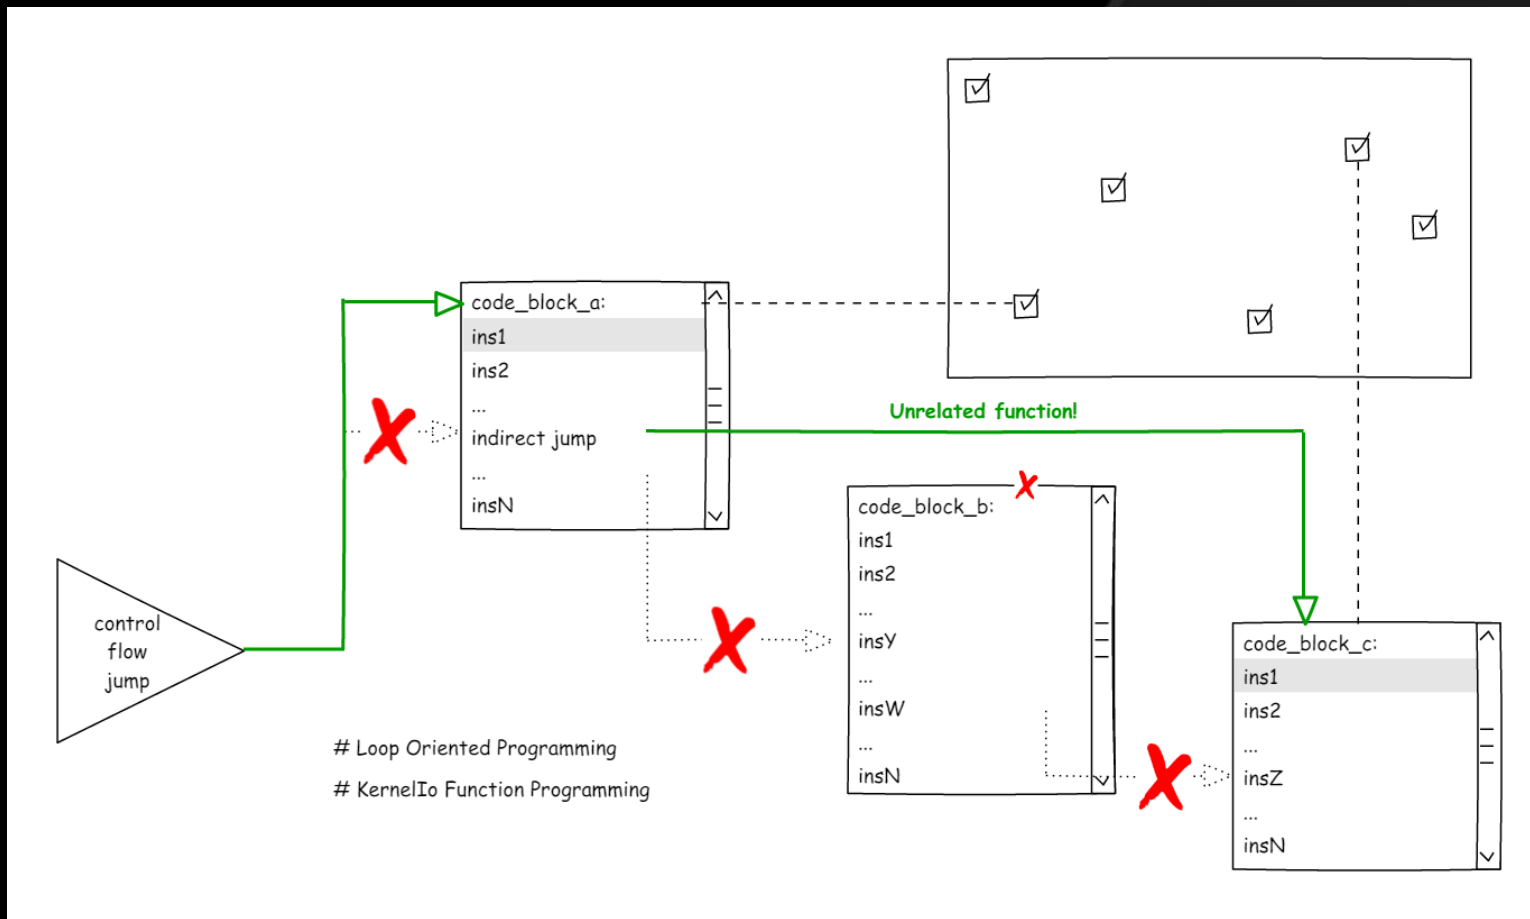
{C/R}FG vs CET vs CFI

- CET-IBT
 - some criticize that it lacks type integrity hash checks
 - It sets bare metal rules – how calls / ret must behave
 - CET is **not** about control flow **integrity** in particular
 - **Implies : ROP is over but you can jump to **any** function instead!**
- CFG + RFG
 - very close to CET
 - Added some integrity checks on top of it
 - black-list (not include into bitmap) “dangerous” functions
 - **Implies : ROP is over but you can jump to large **subset** of functions instead!**
 - depends on corner cases however
 - As you can see CFG improvements over time, cat – mouse game ?

ROP – wild technique



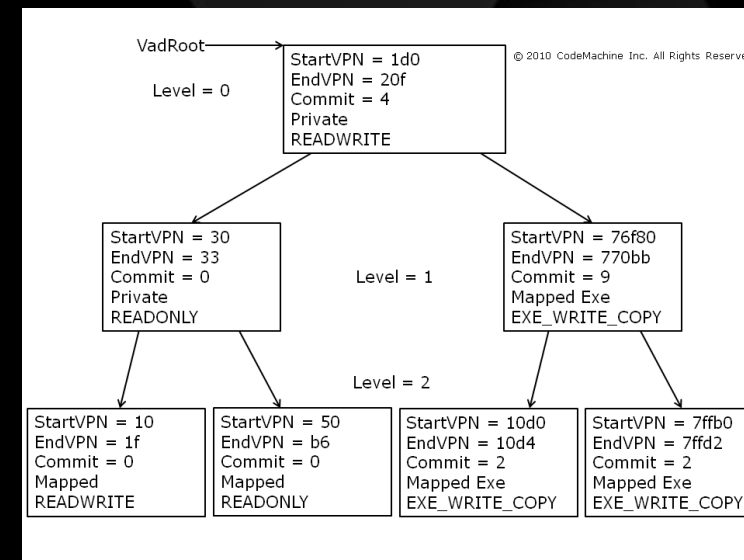
Function oriented - restricted technique



CET-{IBT+ShadowStack} / {C/R}FG

- Gadgets == functions (Loop dispatcher / Kernello)
 - Proper CFI needed (clang / pax)
- Kernel Code Exec, doable ?
 - Well, if you want install old school rootkit, doable
 - RWE on kernel component in your reach
 - Patch legitimate kernel
 - Jump
 - Patch back
 - Obviously bad code!
 - any reasonable benefits from kernel code exec ?
- Kernel Code Exec, needed ?
 - do you even need your own _eproc ?
 - why don't just use kernel io (legit non-harm) syscalls
 - force(trick) other _ethreads to do job for you
 - Patching their stacks (params / rets) should do the job

(VadRoot technique)



<http://www.pcc.me.uk/~peter/acad/usenix14.pdf> <http://www.slideshare.net/PeterHlavaty/back-to-the-core>

<https://pax.grsecurity.net/docs/PaXTeam-H2HC15-RAP-RIP-ROP.pdf>

Data attacks! - TODO

What is your goal ?

DOP – Turing Complete

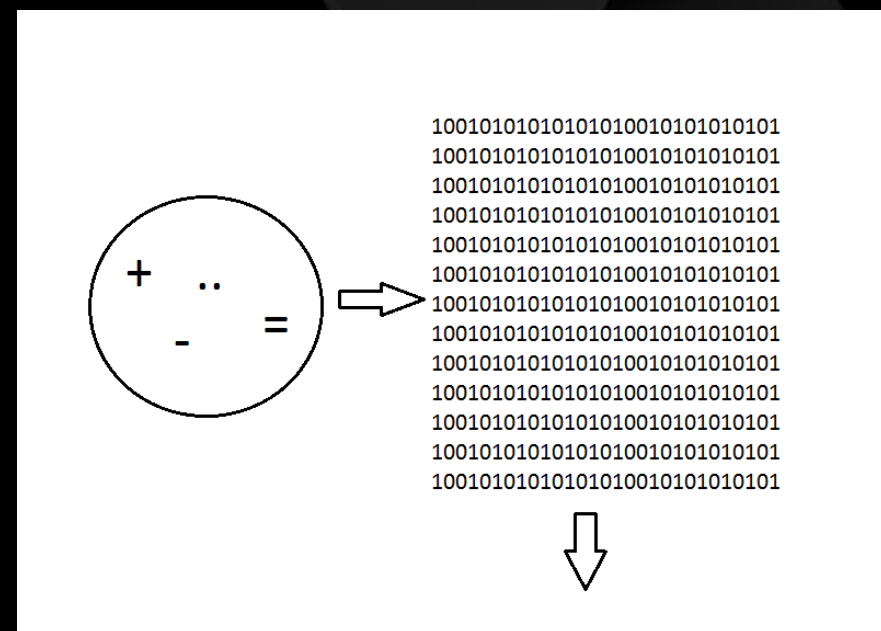
```
1 //dispatcher & jump :
2 void cmd_loop(server_rec *server,conn_t *c) {
3     while (TRUE) {
4         pr_netio_telnet_gets(buf, ...);
5         cmd = make_ftp_cmd(buf, ...);
6         pr_cmd_dispatch(cmd); //calls functions
7                               // with memory errors and gadgets
8     }
9 }
10 char *pr_netio_telnet_gets(char *buf,...) {
11     while(*pbuf->current != '\n' && toread>0)
12         //reads through virtual PC
13         *buf++ = *pbuf->current++;
14 }
```

Code 10. Gadget dispatcher and simulated jump gadget. `pbuf->current` is the virtual PC pointing to the malicious input.

- (Conditional) jump operation. Code 10 shows the ProFTPD program logic to read the next command from an input buffer. `pbuf->current` is a pointer to the next command in the input, thus forming a virtual PC for the attacker's MINDOP program. By corrupting `pbuf->current`, the attacker can select a particular input that invokes a specific MINDOP operation. We use the assignment operation to conditionally update the virtual PC, thus simulating a conditional jump operation.

Kernel io


- In the end data are all what matters!
- Code is just group of instructions working over those data
- Your ultimate goal is to access data
- No need to be SYSTEM
- No need to do Kernel Exec
 - Kernel io can 'emulate' apis
 - You can do math operations
 - ...



Data attack prevention

<https://taesoo.gtisc.gatech.edu/pubs/2016/song:kenali.pdf>

- Same piece of code should work on same memory type
- w32k gdi function will not touch _eprocess
- possible to simple set ranges (hardware) ?
- Pool isolation
- Isolated heap (object caches) ?
- local vars (thread) by default allowed
- Existing solutions ?
 - Intel mpx – seems not going to succeed ?
 - PAC – looks like game changer (CPI approach) ?
 - Isolation ? 😊

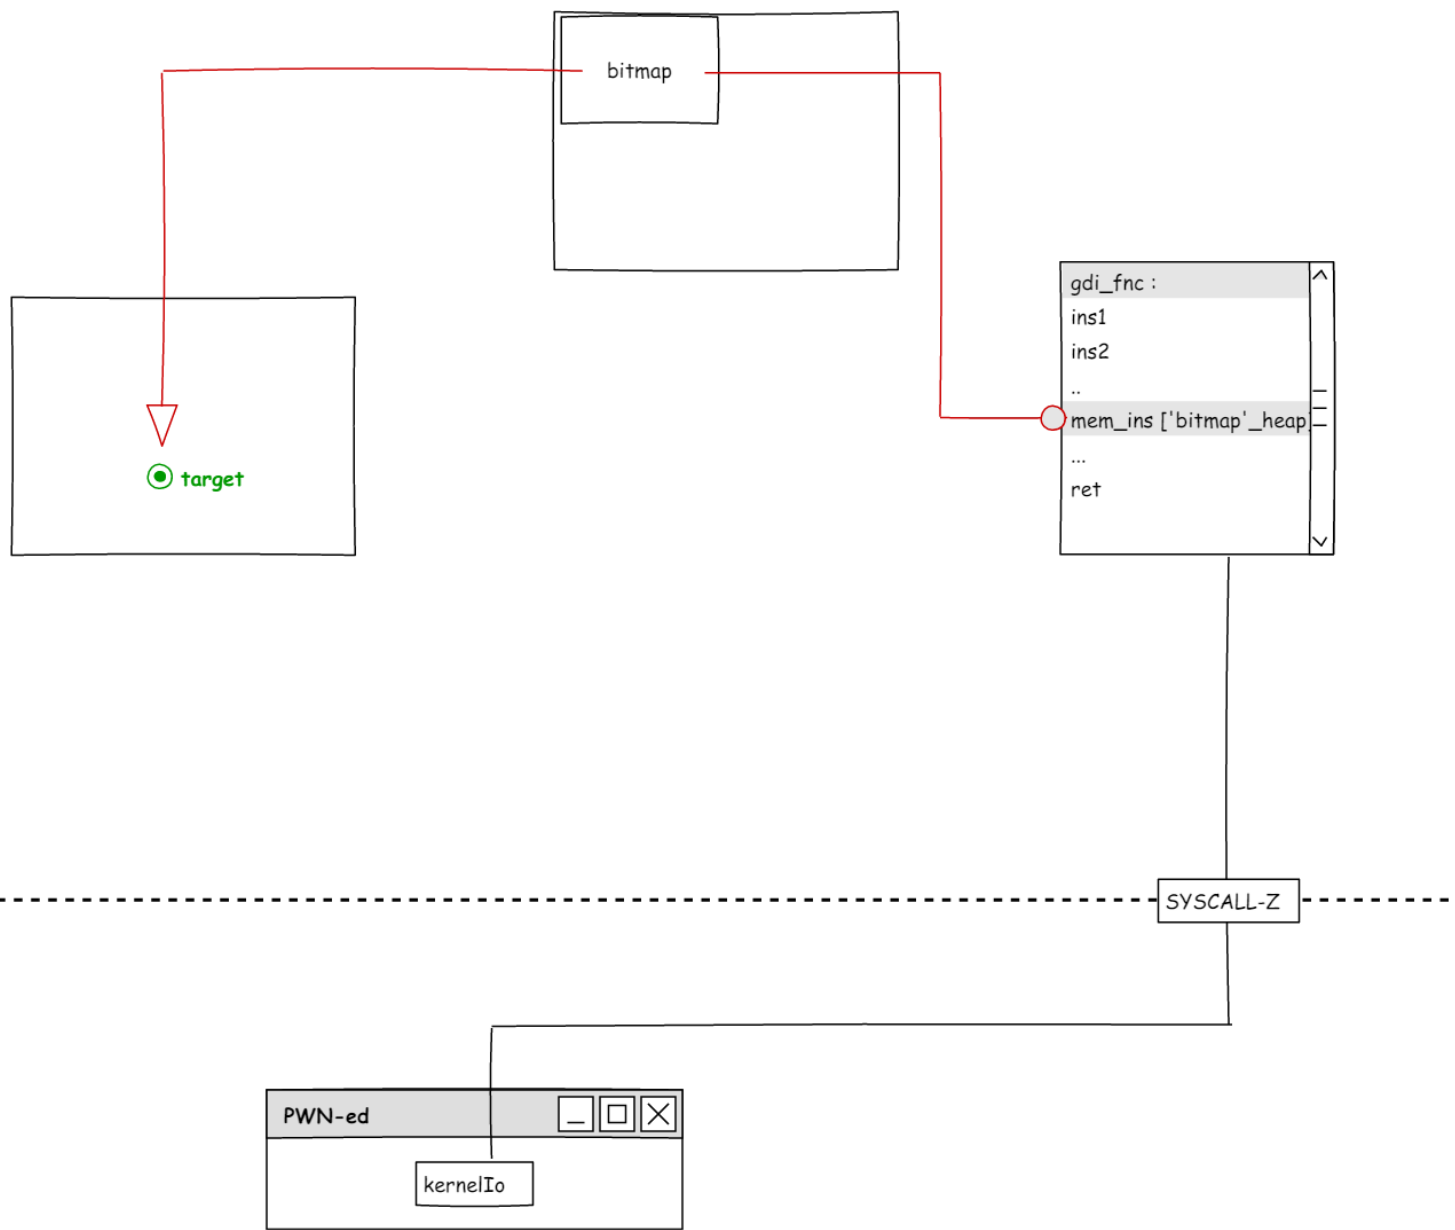


```
kernel_fnc_working_over_mem(  
    __inout type_t* mem  
)  
{  
    ...  
    //setting mem access limits  
    SETMB typeinfo<type_t>.base  
    SETML typeinfo<type_t>.limit  
    //access memory  
    mov [mem], reg  
    mov reg, [mem]  
    ...  
}
```

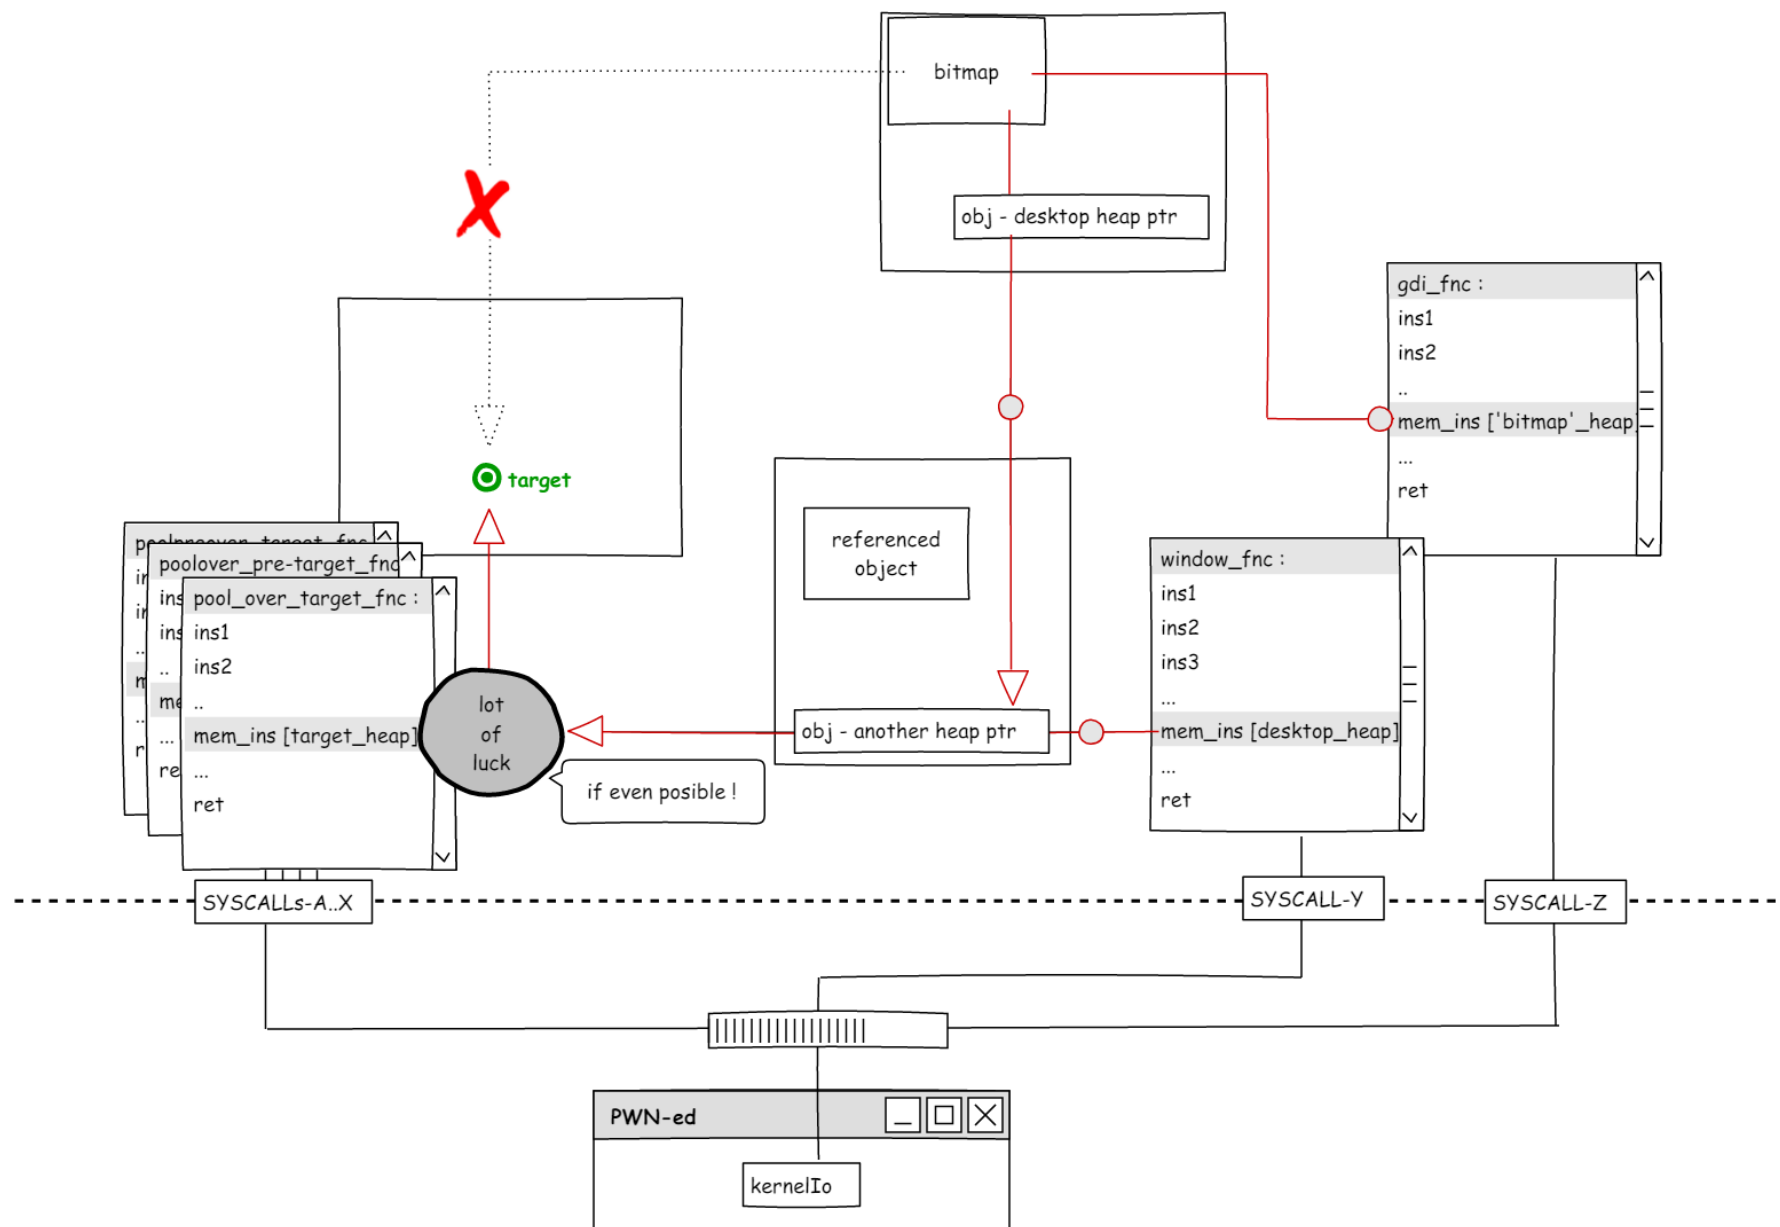
<https://software.intel.com/en-us/blogs/2013/07/22/intel-memory-protection-extensions-intel-mpx-support-in-the-gnu-toolchain>

<https://community.arm.com/groups/processors/blog/2016/10/27/armv8-a-architecture-2016-additions>

Kernel Io vs cross heap checks



Kernel Io vs cross heap checks



Application Guard!

- well, this can change things lot!
- still kernel-io is here real power ?
 - ... well ...
 - doomed within domain
- One does not simple do exec when ring -1 is watching!
 - no real RWE page [ept]
 - code signing [rw -> e]

Containers

- But, as always .. it is just another layer, another ring ...
 - bugs
 - logical issues
- with small big difference
 - less code
 - code quality ++

conclusions

- KASLR, nullptr-deref, Nx, SMEP
 - very well adapted!
- ACL for kernel object access
 - good security model!
- w32k re-designing { filtering & lockdown }
 - tackling biggest security problem in windows kernel!
- {C/R}FG
 - nice anti-exploitation approaches! { - step by step to adapt to kernel }
- CET - IBT & Shadow Stack
 - nice co-operation towards to new security model!
- HyperV based App Containers
 - moving operating system to another league in terms of security

conclusions - todo#memory-corruptions

- Control flow integrity – to improve
 - fine grained CFG – fast, smart, compat! but can it be fine enough ?
 - adapt { forwarded-edge (clang) / type-hash (pax) } – existing, fast, does the job!
 - Make code reuse attacks to disappear ?
- Data attacks break down – to tackle
 - Memory access boundary checks
 - pointer load / function to heap memory access
 - Make kernel io disappear ?
- Will it move memory corruptions from attack vector to **state of art** ?
 - memory corruption == ddos only ?
 - Not yet, long way to go, but now we can see that possibility
 - However maybe that can even happen in foreseeable future ?



Thank you!
Q & A

科恩实验室
KEEN
security lab

Tencent

